

# Large-Scale Machine Learning (1)

CS246: Mining Massive Datasets  
Jure Leskovec, Stanford University  
<http://cs246.stanford.edu>



# New Topic: ML!

## High dim. data

Locality  
sensitive  
hashing

Clustering

Dimensional  
ity  
reduction

## Graph data

PageRank,  
SimRank

Community  
Detection

Spam  
Detection

## Infinite data

Filtering  
data  
streams

Web  
advertising

Queries on  
streams

## Machine learning

Decision  
Trees

Neural  
Networks

Parallel SGD

## Apps

Recommen  
der systems

Association  
Rules

Duplicate  
document  
detection

# Supervised Learning

**Given some data:**

- “Learn” a function to map from the **input** to the **output**
- **Given:**  
Training examples  $(x_i, y_i = f(x_i))$  for some unknown function  $f$
- **Find:**  
A good approximation to  $f$

# Many Other ML Paradigms

- **Supervised:**

- Given “labeled data”  $\{x, y\}$ , learn  $f(x) = y$

- **Unsupervised:**

- Given only “unlabeled data”  $\{x\}$ , learn  $f(x)$

- **Semi-supervised:**

- Given some labeled  $\{x, y\}$  and some unlabeled data  $\{x\}$ , learn  $f(x) = y$

- **Active learning:**

- When we predict  $f(x) = y$ , we then receive true  $y^*$

- **Transfer learning:**

- Learn  $f(x)$  so that it works well on new domain  $f(z)$

# Supervised Learning

- Would like to do **prediction**:  
estimate a function  $f(x)$  so that  $y = f(x)$
- Where **y** can be:
  - **Real number**: Regression
  - **Categorical**: Classification
  - **Complex object**:
    - Ranking of items, Parse tree, etc.
- Data is **labeled**:
  - Have many pairs  $\{(x, y)\}$ 
    - $x$  ... vector of binary, categorical, real valued features
    - $y$  ... class, or a real number

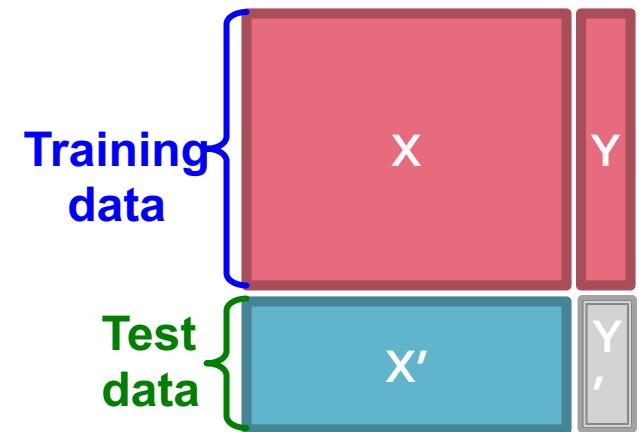
# Supervised Learning

- **Task:** Given data  $(X, Y)$  build a model  $f()$  to predict  $Y'$  based on  $X'$

- **Strategy:** Estimate  $y = f(x)$  on  $(X, Y)$ .

Hope that the same  $f(x)$  also works to predict unknown  $Y'$

- The “hope” is called **generalization**
  - **Overfitting:** If  $f(x)$  predicts well  $Y$  but is unable to predict  $Y'$
- **We want to build a model that generalizes well to unseen data**



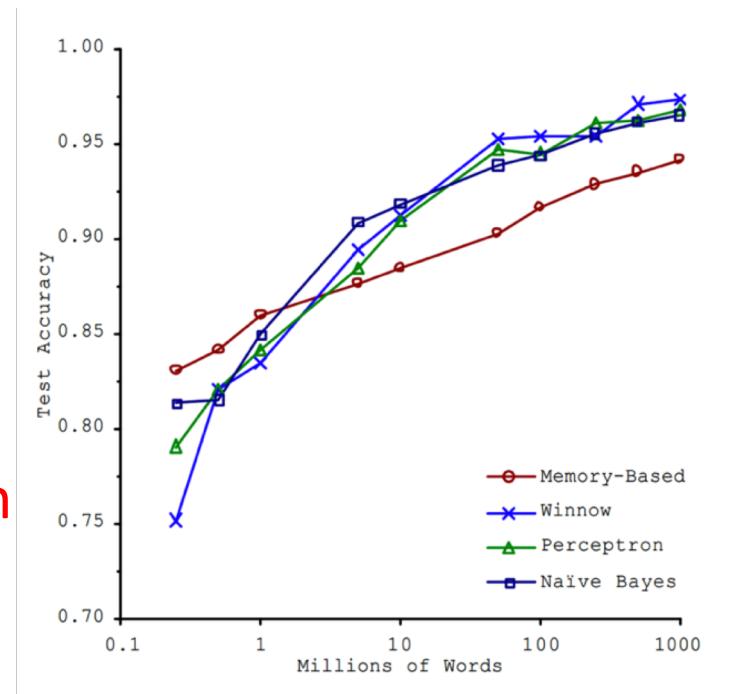
# Why Large-Scale ML?

## ■ Brawn or Brains?

- In 2001, Microsoft researchers ran a test to evaluate 4 of different approaches to ML translation

## ■ Findings:

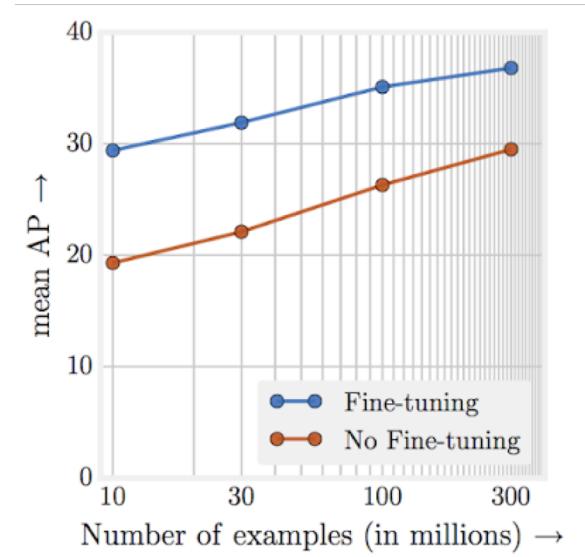
- Size of the dataset used to train the model mattered more than the model itself
- As the dataset grew large, performance difference between the models became small



Banko, M. and Brill, E. (2001), "[Scaling to Very Very Large Corpora for Natural Language Disambiguation](#)"

# Why Large-Scale ML?

- **The Unreasonable Effectiveness of Data**
  - In 2017, Google revisited the same type of experiment with the latest Deep Learning models in computer vision
- **Findings:**
  - Performance increases logarithmically based on volume of training data
  - Complexity of modern ML models (i.e., deep neural nets) allows for even further performance gains
- **Large datasets + large ML models => amazing results!!**



"Revisiting Unreasonable Effectiveness of Data in Deep Learning Era": <https://arxiv.org/abs/1707.02968>

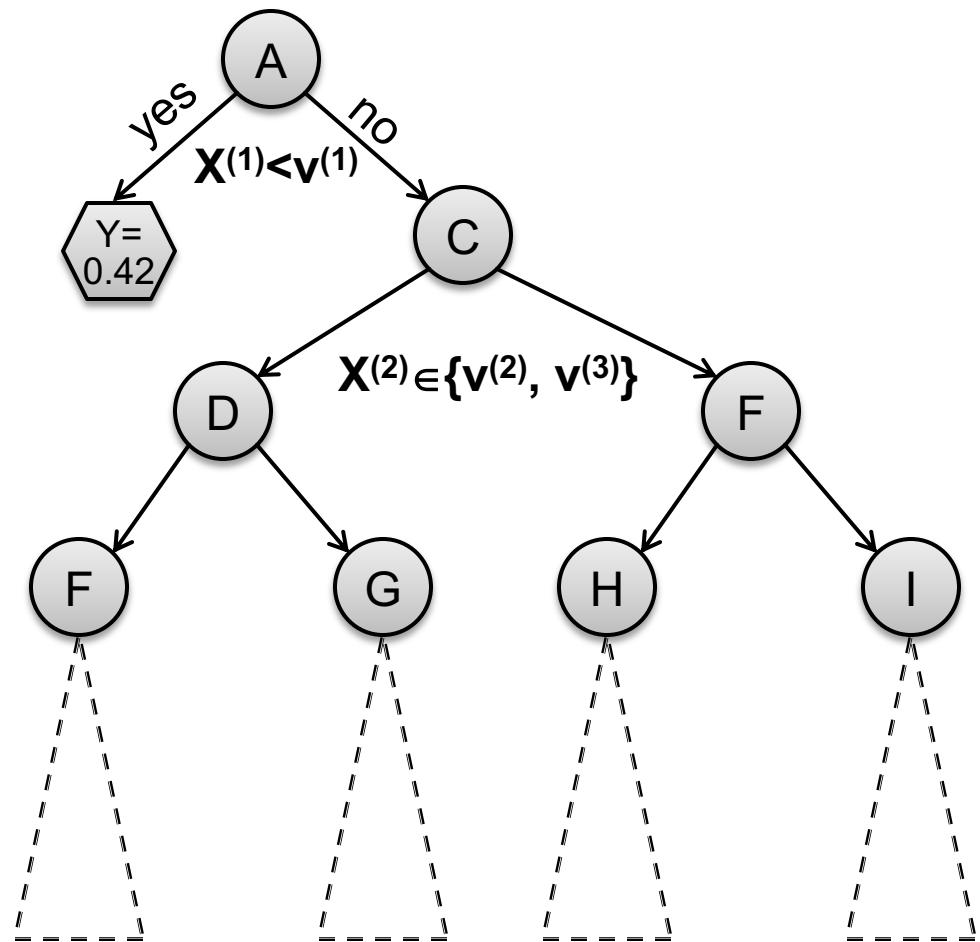
# Decision Trees

# Decision Tree Learning

- Given one attribute (e.g., lifespan), try to predict the value of new people's lifespans by means of some of the other available attribute
- Input attributes:
  - $d$  features/attributes:  $x^{(1)}, x^{(2)}, \dots, x^{(d)}$
  - Each  $x^{(j)}$  has domain  $O_j$ 
    - Categorical:  $O_j = \{red, blue\}$
    - Numerical:  $H_j = (0, 10)$
  - $Y$  is output variable with domain  $O_Y$ :
    - Categorical: Classification, Numerical: Regression
- Data D:
  - $n$  examples  $(x_i, y_i)$  where  $x_i$  is a  $d$ -dim feature vector,  $y_i \in O_Y$  is output variable
- Task:
  - Given an input data vector  $x$  predict  $y$

# Decision Trees

- A **Decision Tree** is a tree-structured plan of a set of attributes to test in order to predict the output



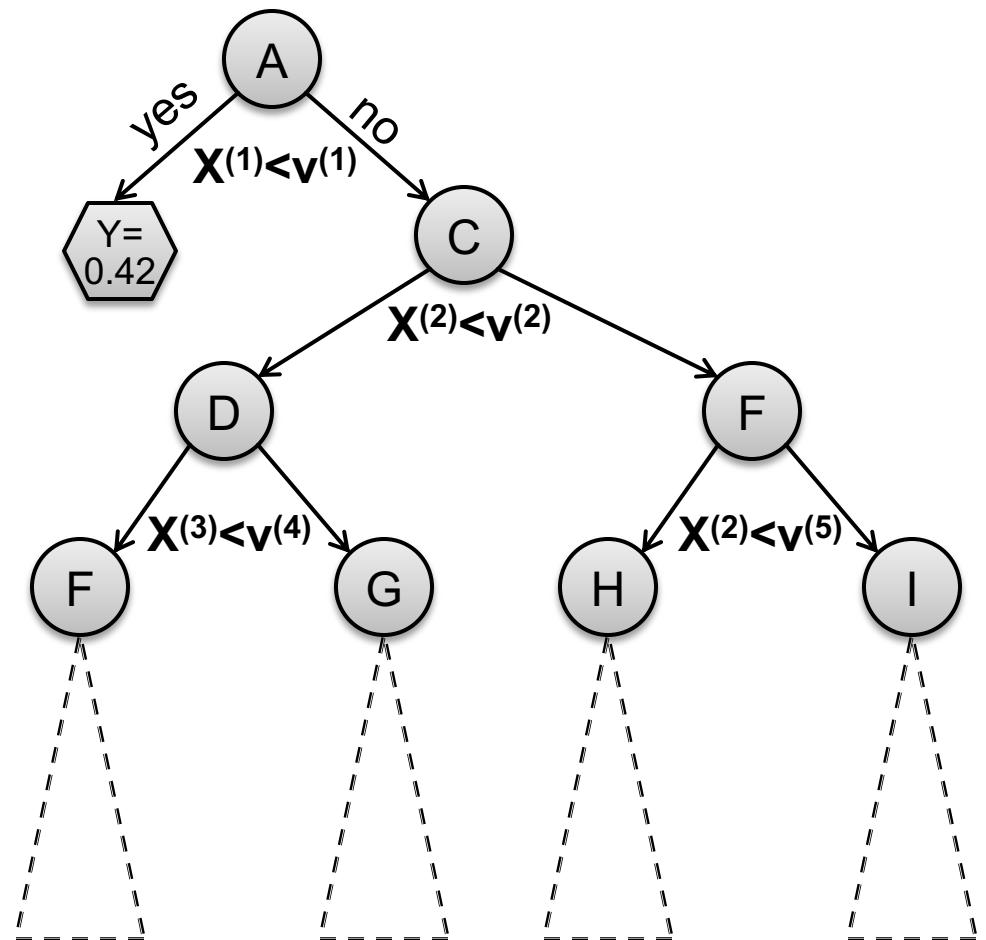
# Decision Trees

- **Decision trees:**

- Split the data at each internal node
- Each leaf node makes a prediction

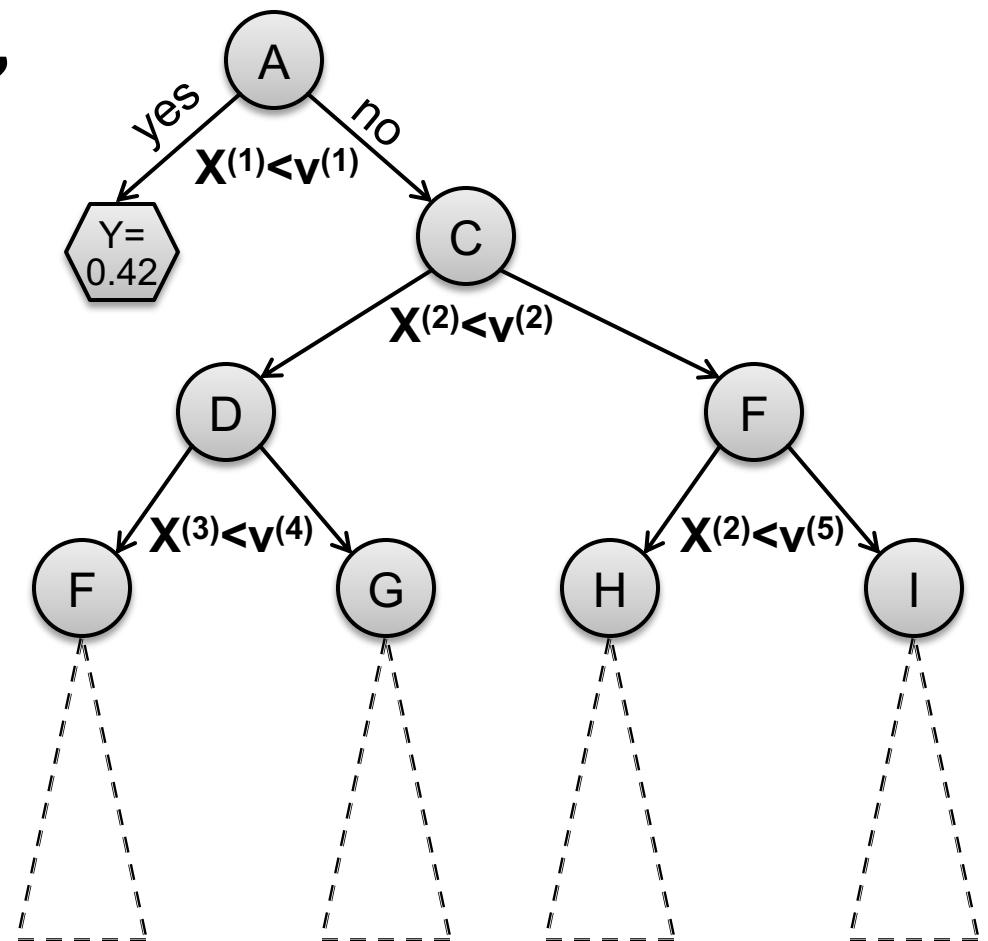
- **Lecture today:**

- Binary splits:  $X^{(j)} < v$
- Numerical attributes
- Regression



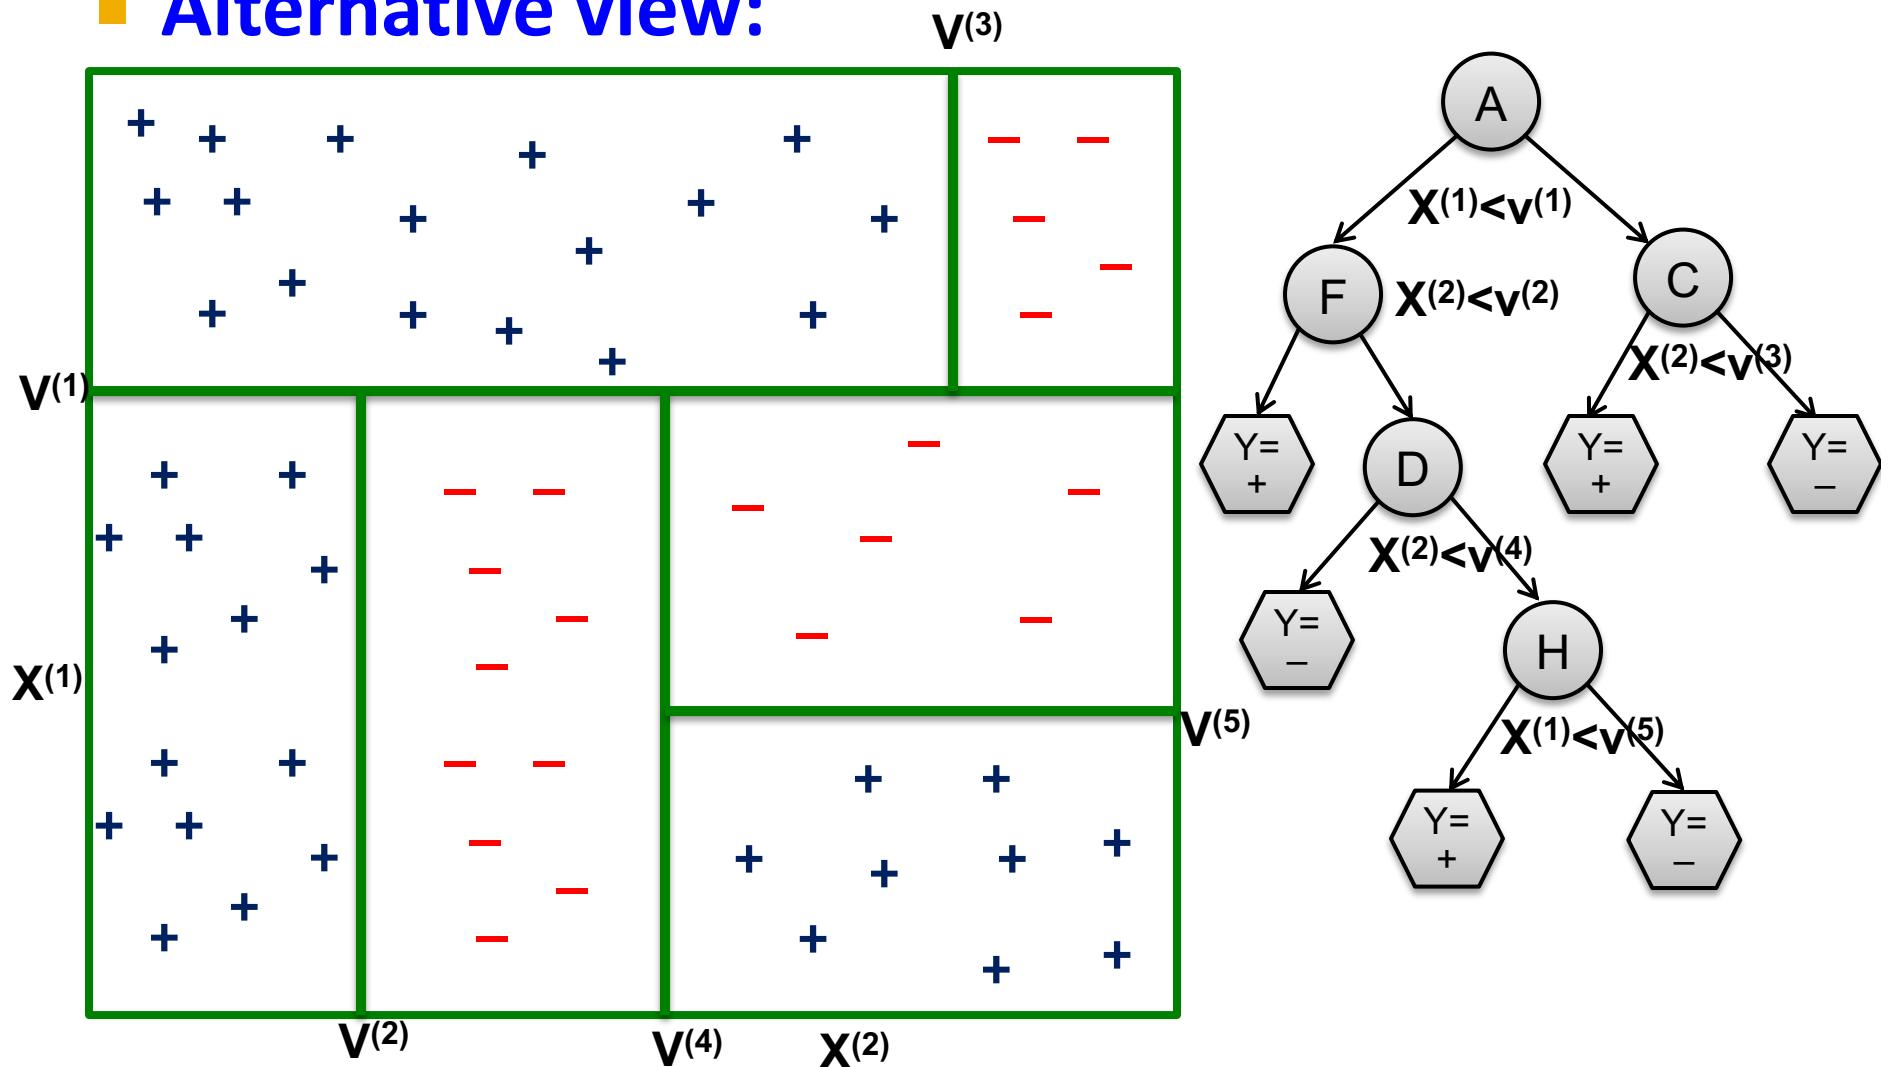
# How to make predictions?

- **Input:** Example  $x_i$
- **Output:** Predicted  $y_i'$
- “Drop”  $x_i$  down the tree until it hits a leaf node
- Predict the value stored in the leaf that  $x_i$  hits



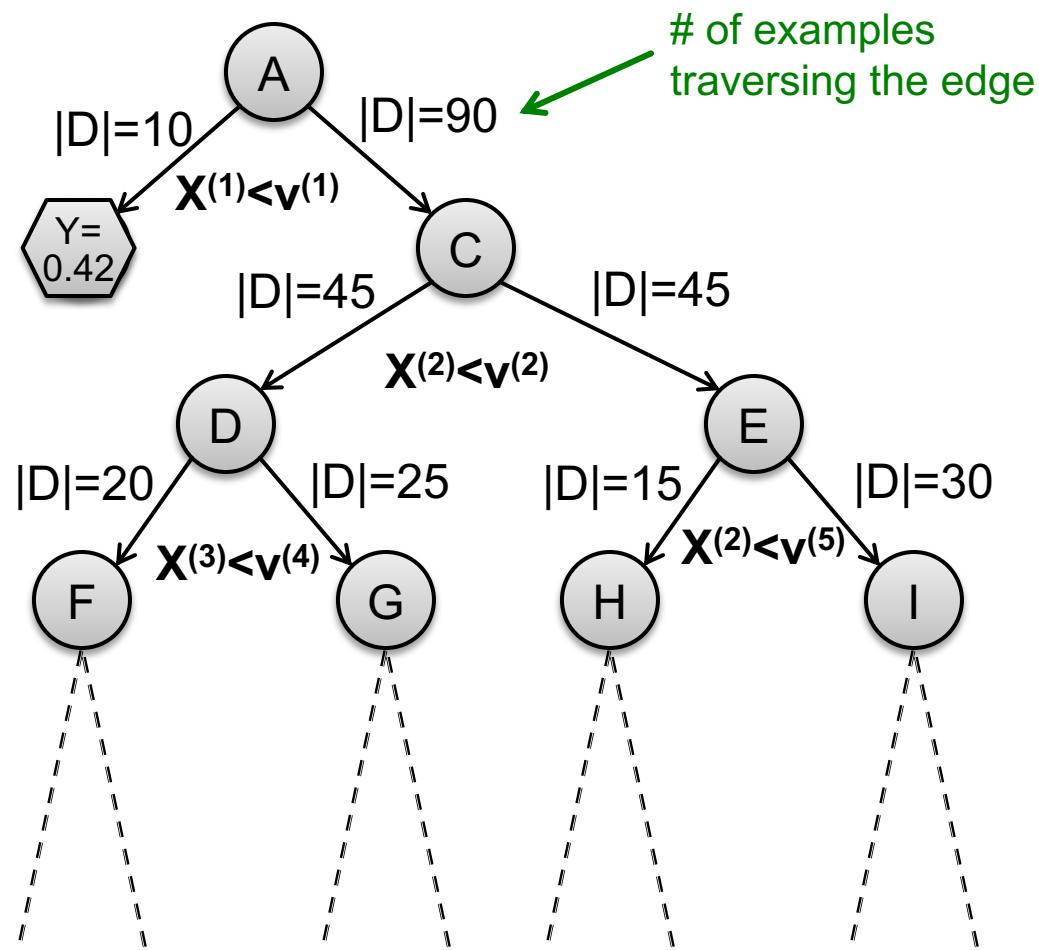
# Decision Trees: feature space

## ■ Alternative view:



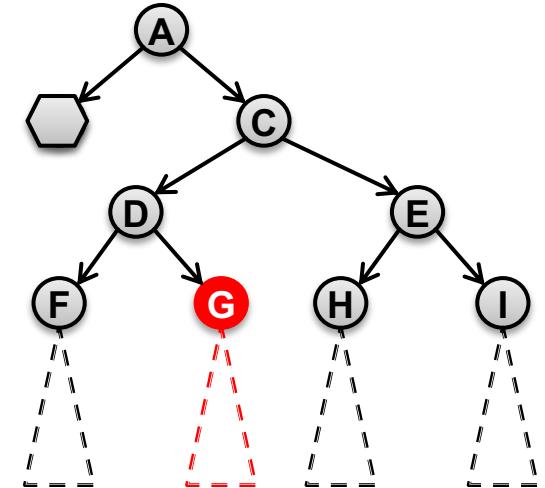
# How to construct a tree?

- Training dataset  $D^*$ ,  $|D^*| = 100$  examples



# How to construct a tree?

- Imagine we are currently at some node  $G$ 
  - Let  $D_G$  be the data that reaches  $G$
- **There is a decision we have to make: Do we continue building the tree?**
  - **If yes**, which variable and which value do we use for a **split**?
    - Continue building the tree recursively
  - **If not**, how do we make a prediction?
    - We need to build a “**predictor node**”



# 3 steps in constructing a tree

---

## Algorithm 1

### BuildSubtree

Require: Node  $n$ , Data  $D \subseteq D^*$

1:  $(n \rightarrow \text{split}, D_L, D_R) = \text{FindBestSplit}(D)$  (1)

2: if  $\text{StoppingCriteria}(D_L)$  then (2)

3:    $n \rightarrow \text{left\_prediction} = \text{FindPrediction}(D_L)$  (3)

4: else

5:           BuildSubtree  $(n \rightarrow \text{left}, D_L)$

6: if  $\text{StoppingCriteria}(D_R)$  then

7:    $n \rightarrow \text{right\_prediction} = \text{FindPrediction}(D_R)$

8: else

9:           BuildSubtree  $(n \rightarrow \text{right}, D_R)$

---

- Requires at least a single pass over the data!

# How to construct a tree?

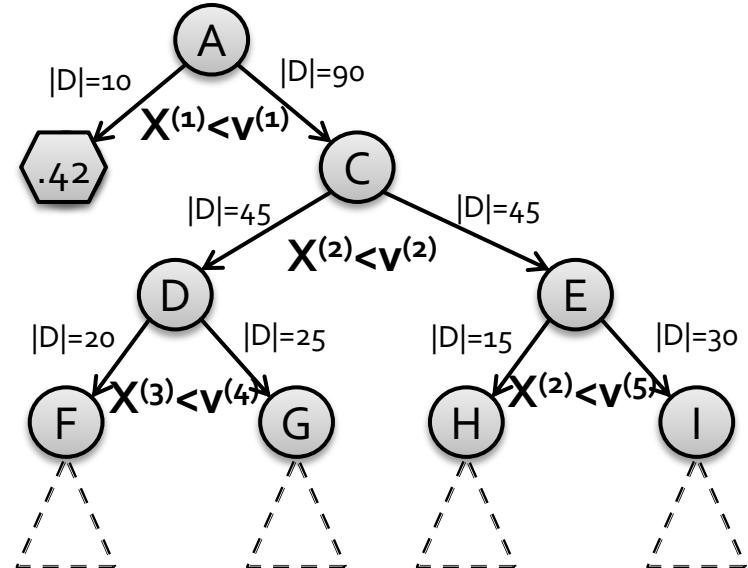
**(1) How to split?** Pick attribute & value that optimizes some criterion

- **Regression: Purity**

- Find split  $(X^{(i)}, v)$  that creates  $D, DL, DR$ : parent, left, right child datasets and **maximizes**:

$$|D| \cdot Var(D) - (|D_L| \cdot Var(D_L) + |D_R| \cdot Var(D_R))$$

- $Var(D) = \frac{1}{n} \sum_{i \in D} (y_i - \bar{y})^2$  ... variance of  $y_i$  in  $D$

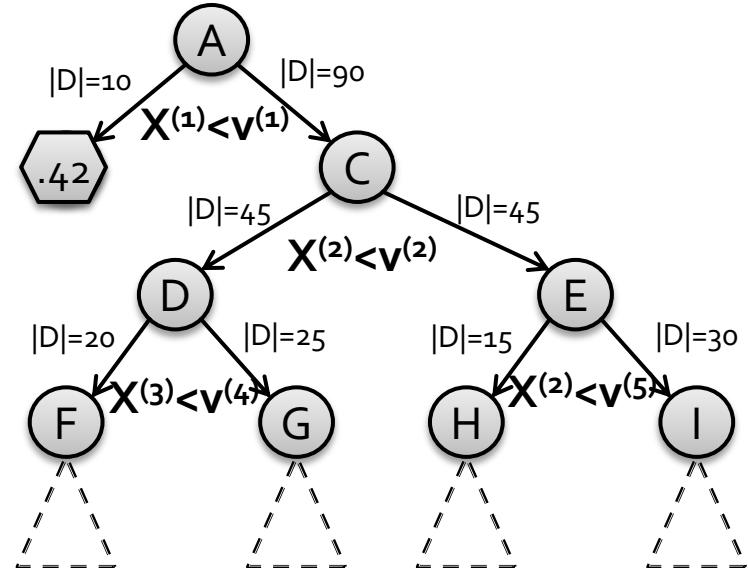


# How to construct a tree?

(1) How to split? Pick attribute & value that optimizes some criterion

- Classification:  
Information Gain

- Measures how much a given attribute  $X$  tells us about the class  $Y$
- $IG(Y | X)$  : We must transmit  $Y$  over a binary link. How many bits on average would it save us if both ends of the line knew  $X$ ?

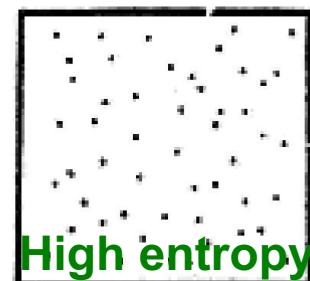
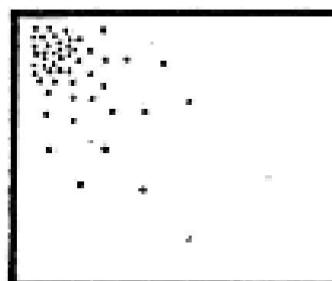
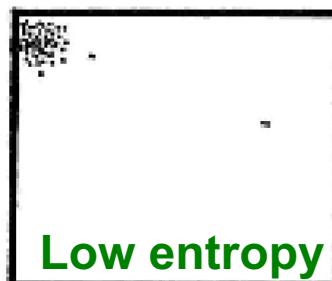


# Why Information Gain? Entropy

**Entropy:** What's the smallest possible number of bits, on average, per symbol, needed to transmit a stream of symbols drawn from  $X$ 's distribution?

**The entropy of  $X$ :**  $H(X) = - \sum_{j=1}^m p(X_j) \log p(X_j)$

- “**High Entropy**”:  $X$  is from a uniform (boring) distribution
  - A histogram of the frequency distribution of values of  $X$  is flat
- “**Low Entropy**”:  $X$  is from a varied (peaks/valleys) distrib.
  - A histogram of the frequency distribution of values of  $X$  would have many lows and one or two highs



# Why Information Gain? Entropy

- Suppose I want to predict  $Y$  and I have input  $X$ 
  - $X = \text{College Major}$
  - $Y = \text{Likes "Casablanca"}$

$X$	$Y$
Math	Yes
History	No
CS	Yes
Math	No
Math	No
CS	Yes
Math	Yes
History	No

- From this data we estimate
  - $P(Y = \text{Yes}) = 0.5$
  - $P(X = \text{Math} \& Y = \text{No}) = 0.25$
  - $P(X = \text{Math}) = 0.5$
  - $P(Y = \text{Yes} | X = \text{History}) = 0$
- Note:
  - $H(Y) = -\frac{1}{2}\log_2(\frac{1}{2}) - \frac{1}{2}\log_2(\frac{1}{2}) = 1$
  - $H(X) = 1.5$

# Why Information Gain? Entropy

- Suppose I want to predict  $Y$  and I have input  $X$ 
  - $X$  = College Major
  - $Y$  = Likes “Casablanca”

X	Y
Math	Yes
History	No
CS	Yes
Math	No
Math	No
CS	Yes
Math	Yes
History	No

- Def: Specific Conditional Entropy
  - $H(Y | X = v)$  = The entropy of  $Y$  among only those records in which  $X$  has value  $v$
  - Example:
    - $H(Y|X = \text{Math}) = 1$
    - $H(Y|X = \text{History}) = 0$
    - $H(Y|X = \text{CS}) = 0$

# Why Information Gain?

- Suppose I want to predict  $Y$  and I have input  $X$ 
  - $X$  = College Major
  - $Y$  = Likes “Casablanca”

$X$	$Y$
Math	Yes
History	No
CS	Yes
Math	No
Math	No
CS	Yes
Math	Yes
History	No

## ■ Def: Conditional Entropy

- $H(Y | X)$  = The average specific conditional entropy of  $Y$ 
  - = if you choose a record at random what will be the conditional entropy of  $Y$ , conditioned on that row's value of  $X$
  - = Expected number of bits to transmit  $Y$  if both sides will know the value of  $X$
- $= \sum_j P(X = v_j)H(Y|X = v_j)$

# Why Information Gain?

- Suppose I want to predict  $Y$  and I have input  $X$ 
  - $H(Y | X)$  = The average specific conditional entropy of  $Y$

X	Y
Math	Yes
History	No
CS	Yes
Math	No
Math	No
CS	Yes
Math	Yes
History	No

$$= \sum_j P(X = v_j) H(Y|X = v_j)$$

- Example:

$v_j$	$P(X=v_j)$	$H(Y X=v_j)$
Math	0.5	1
History	0.25	0
CS	0.25	0

- So:  $H(Y|X)=0.5*1+0.25*0+0.25*0 = 0.5$

# Why Information Gain?

- Suppose I want to predict  $Y$  and I have input  $X$
- Def: Information Gain

■  $IG(Y|X)$  = I must transmit  $Y$ . How many bits on average would it save me if both ends of the line knew  $X$ ?

$$IG(Y|X) = H(Y) - H(Y | X)$$

- Example:

- $H(Y) = 1$
- $H(Y | X) = 0.5$
- Thus  $IG(Y|X) = 1 - 0.5 = 0.5$

X	Y
Math	Yes
History	No
CS	Yes
Math	No
Math	No
CS	Yes
Math	Yes
History	No

# What is Information Gain used for?

- Suppose you are trying to predict whether someone is going to live past 80 years
- From historical data you might find:
  - $IG(LongLife | HairColor) = 0.01$
  - $IG(LongLife | Smoker) = 0.4$
  - $IG(LongLife | Gender) = 0.25$
  - $IG(LongLife | LastDigitOfSSN) = 0.00001$
- IG tells us how much information about  $Y$  is contained in  $X$ 
  - So attribute  $X$  that has high  $IG(Y|X)$  is a good split!

# 3 steps in constructing a tree

---

## Algorithm 1 **BuildSubtree**

---

Require: Node  $n$ , Data  $D \subseteq D^*$

1:  $(n \rightarrow \text{split}, D_L, D_R) = \text{FindBestSplit}(D)$  (1)

2: if  $\text{StoppingCriteria}(D_L)$  then (2)

3:    $n \rightarrow \text{left\_prediction} = \text{FindPrediction}(D_L)$  (3)

4: else

5:                   **BuildSubtree**  $(n \rightarrow \text{left}, D_L)$

6: if  $\text{StoppingCriteria}(D_R)$  then

7:    $n \rightarrow \text{right\_prediction} = \text{FindPrediction}(D_R)$

8: else

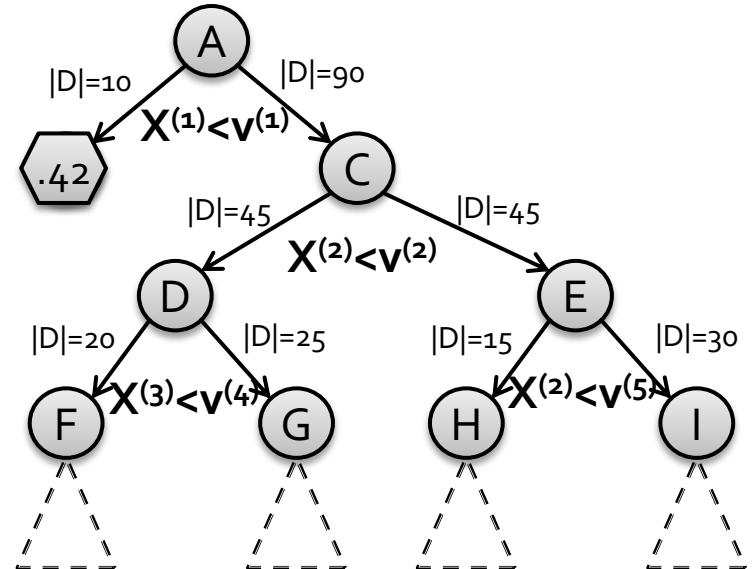
9:                   **BuildSubtree**  $(n \rightarrow \text{right}, D_R)$

---

# When to stop?

## (2) When to stop?

- Many different heuristic options
- Two ideas:
  - (1) When the leaf is “pure”
    - The target variable does not vary too much:  $Var(y_i) < \varepsilon$
  - (2) When # of examples in the leaf is too small
    - For example,  $|D| \leq 100$



# How to predict?

## (3) How to predict?

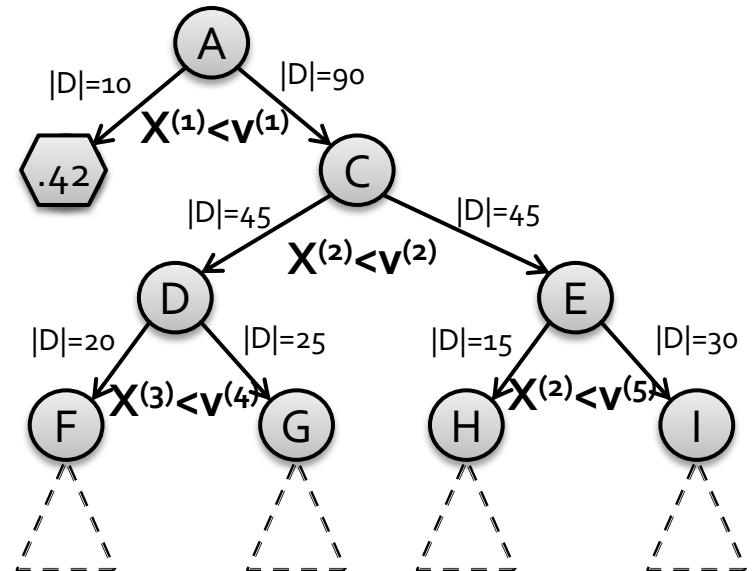
- Many options

- Regression:

- Predict average  $y_i$  of the examples in the leaf
    - Build a linear regression model on the examples in the leaf

- Classification:

- Predict most common  $y_i$  of the examples in the leaf



# **Building Decision Trees Using MapReduce**

# Problem: Building a tree

- Given a large dataset with hundreds of attributes
- Build a decision tree!
- General considerations:
  - Tree is small (can keep it memory):
    - Shallow (~10 levels)
  - Dataset too large to keep in memory
  - Dataset too big to scan over on a single machine
  - MapReduce to the rescue!

---

## Algorithm 1 BuildSubTree

---

```
Require: Node  $n$ , Data  $D \subseteq D^*$ 
1:  $(n \rightarrow \text{split}, D_L, D_R) = \text{FindBestSplit}(D)$ 
2: if StoppingCriteria( $D_L$ ) then
3:    $n \rightarrow \text{left\_prediction} = \text{FindPrediction}(D_L)$ 
4: else
5:   BuildSubTree( $n \rightarrow \text{left}, D_L$ )
6: if StoppingCriteria( $D_R$ ) then
7:    $n \rightarrow \text{right\_prediction} = \text{FindPrediction}(D_R)$ 
8: else
9:   BuildSubTree( $n \rightarrow \text{right}, D_R$ )
```

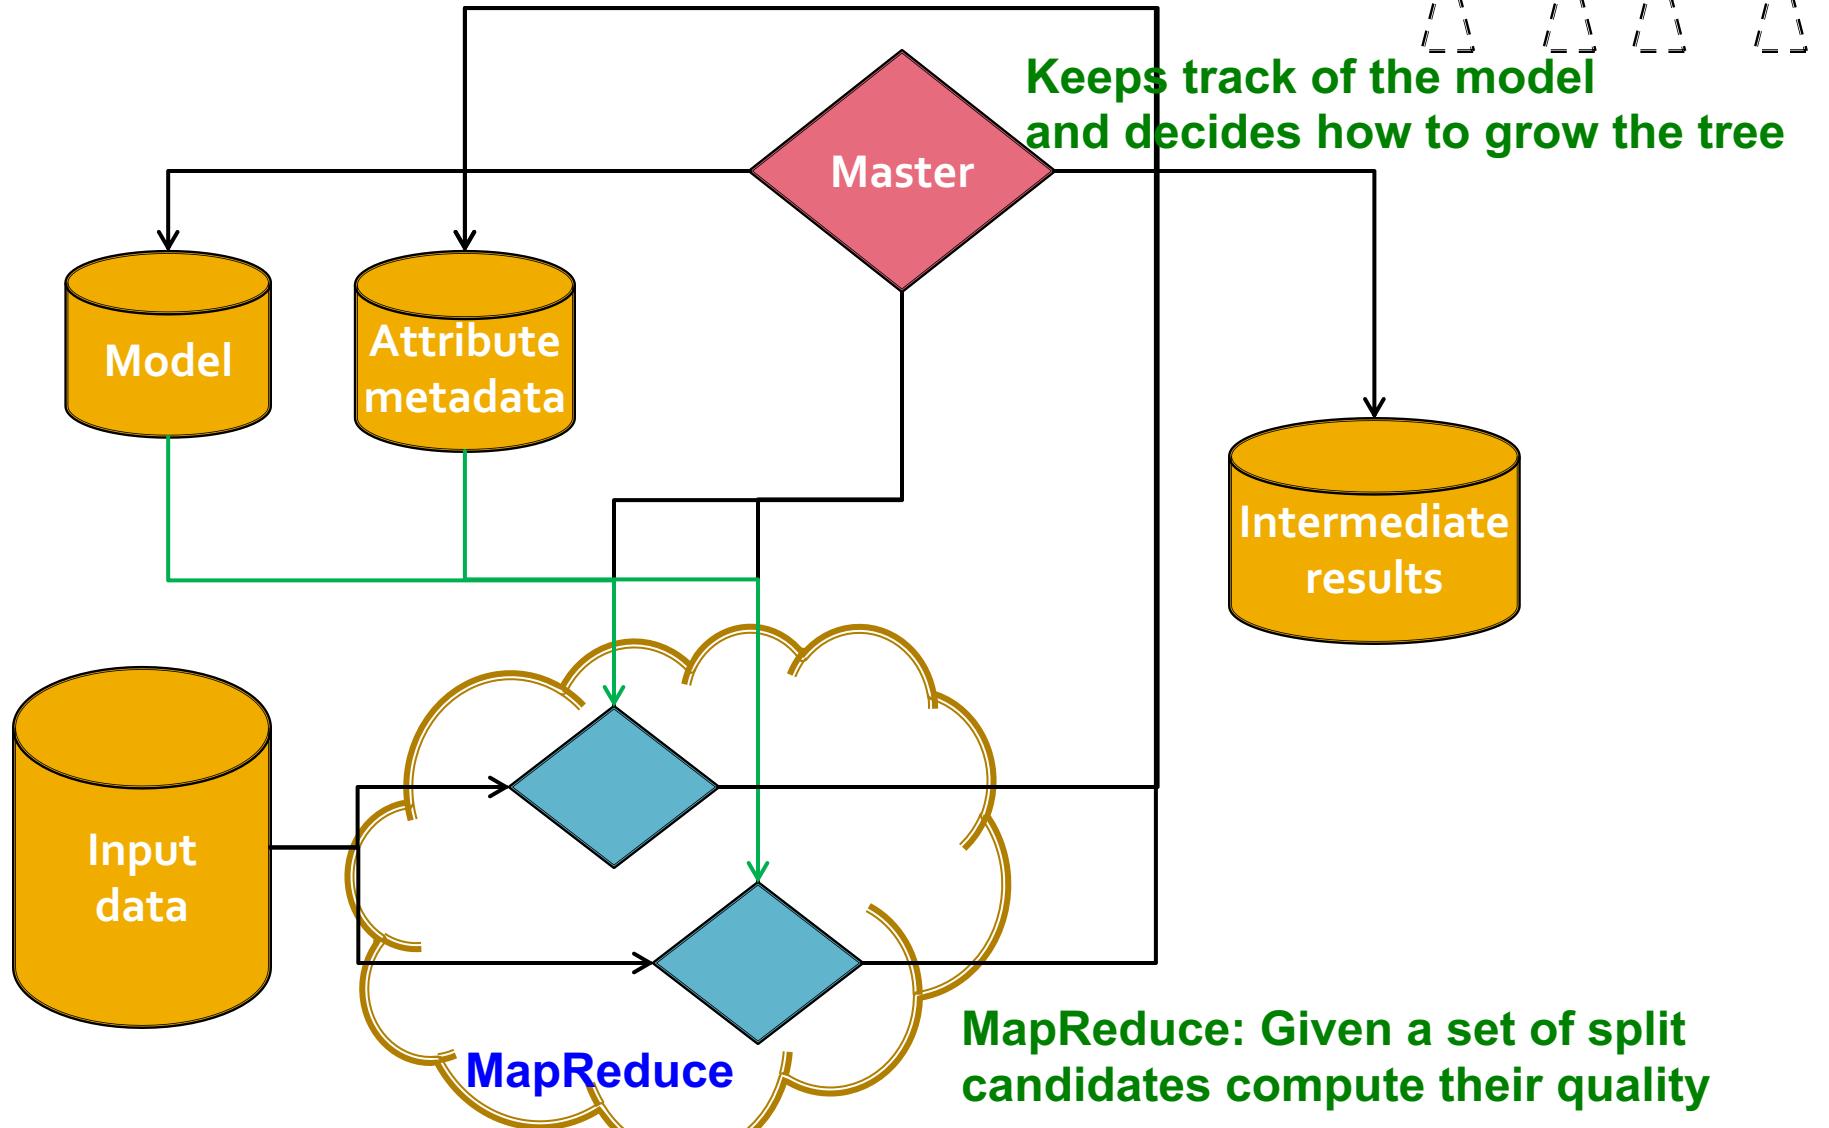
---

# Today's Lecture: PLANET

## Parallel Learner for Assembling Numerous Ensemble Trees [Panda et al., VLDB '09]

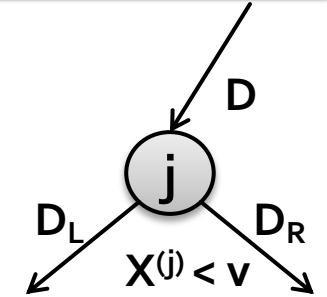
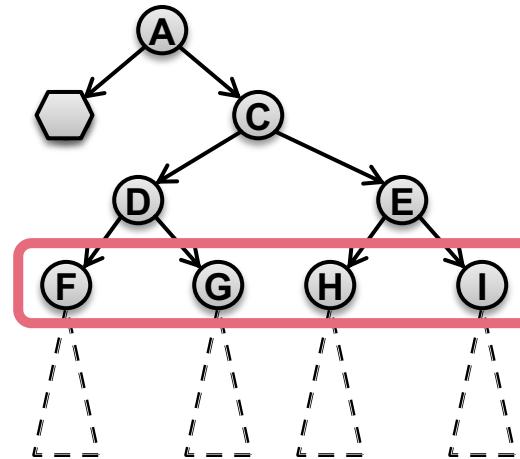
- A **sequence** of MapReduce jobs that builds a decision tree
- Spark MLlib Decision Trees are based on PLANET
- **Setting:**
  - Hundreds of **numerical** (discrete & continuous, but not categorical) attributes
  - Target variable is **numerical**: **Regression**
  - Splits are **binary**:  $X^{(j)} < v$
  - **Decision tree is small enough for each Mapper to keep it in memory**
  - **Data too large to keep in memory**

# PLANET Architecture



# PLANET: Building the Tree

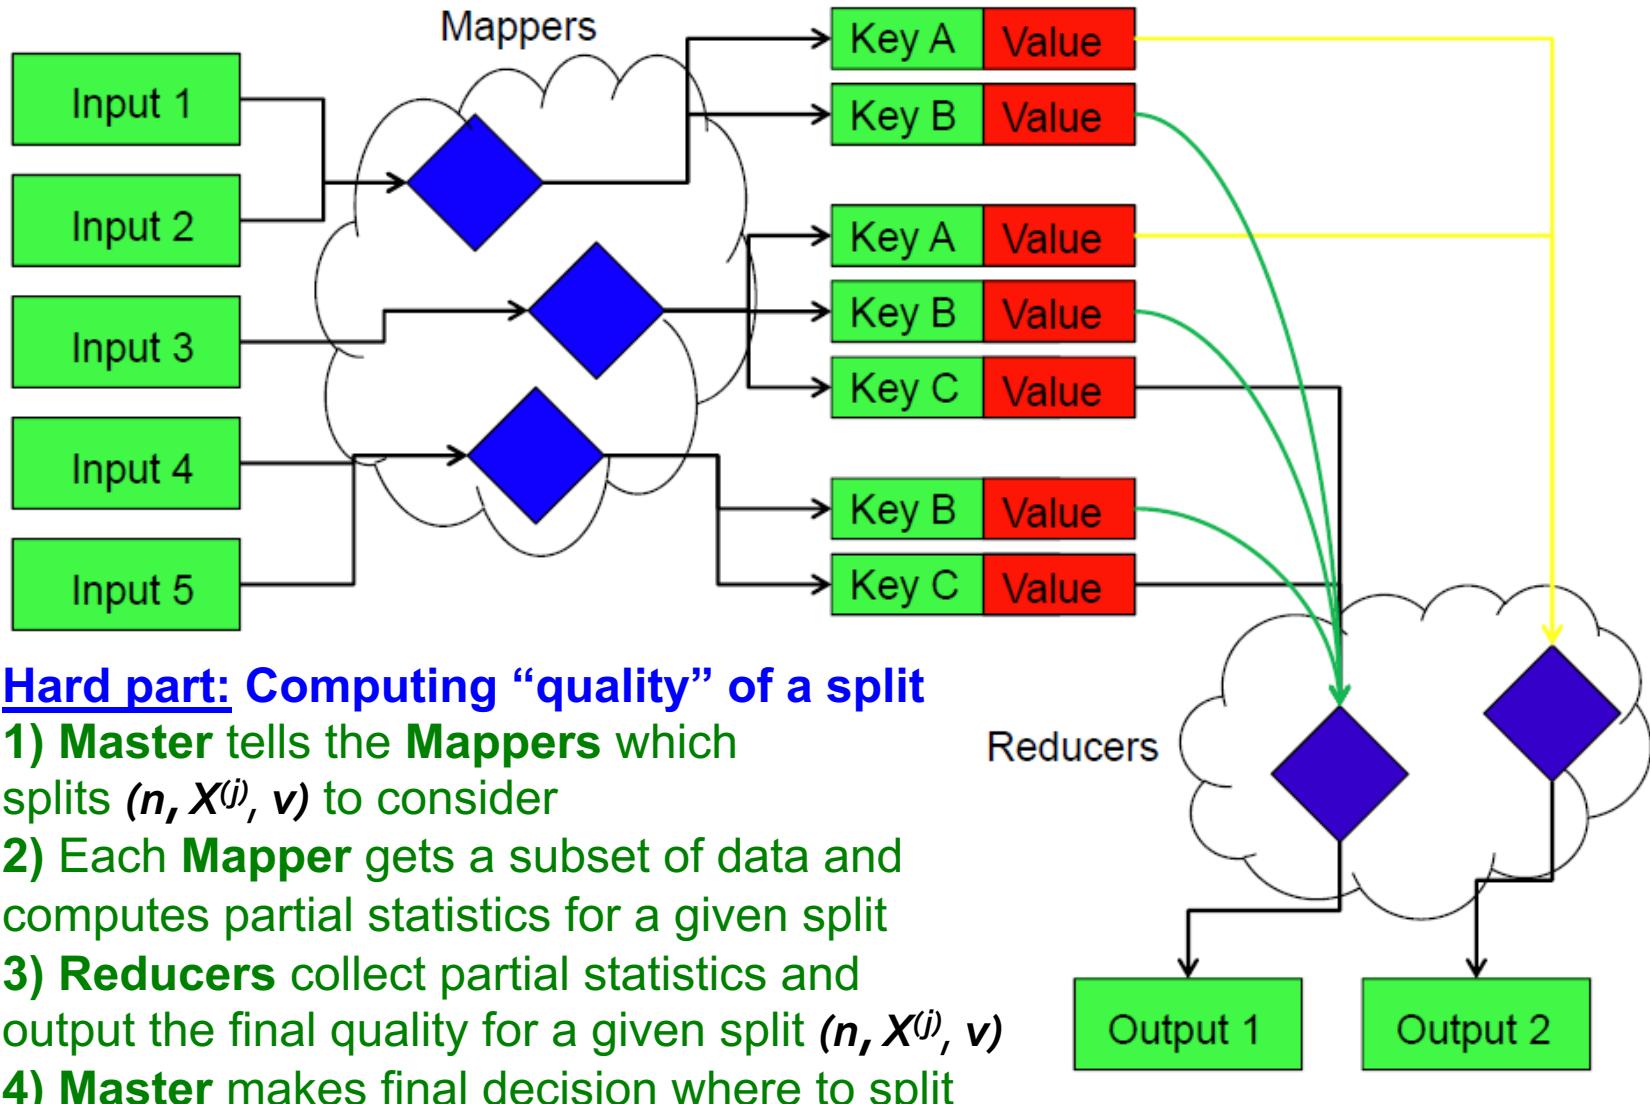
- The tree will be built in levels
  - One level at a time:



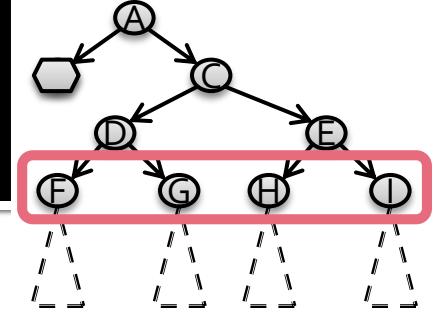
## Steps:

- 1) Master decides candidate splits ( $n, X^{(j)}, v$ )
- 2) MapReduce computes quality of those splits
- 3) Master then grows the tree for a level
- 4) Goto 1)

# Decision trees on MapReduce

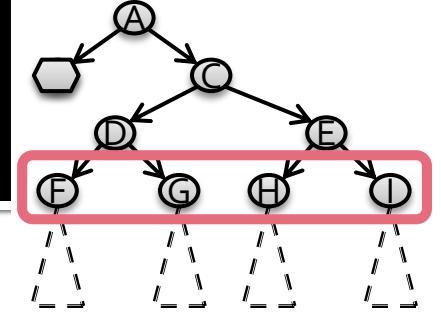


# PLANET Overview



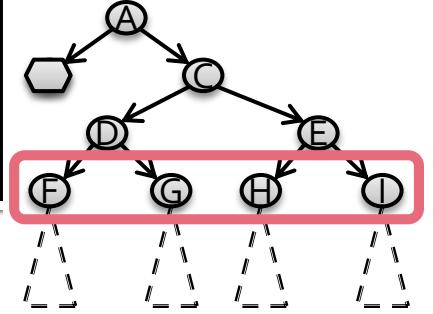
- We build the tree level by level
  - One MapReduce step builds **one level of the tree**
- **Mapper**
  - Considers a number of candidate splits (**node, attribute, value**) on its subset of the data
  - For each split it stores **partial statistics**
  - Partial split-statistics is sent to **Reducers**
- **Reducer**
  - Collects all partial statistics and determines best split
- **Master** grows the tree for one level

# PLANET Overview



- **Mapper** loads the **DT model** and info about which **attribute splits** (split is a tuple  $\langle \text{NodeID}, \text{Attribute}, \text{Value} \rangle$ ) to consider
  - Each mapper sees a subset of the data  $D^*$
  - Mapper “drops”/classifies each datapoint  $d$  using the tree to find the leaf node  $L$  where  $d$  lands
  - For each leaf node  $L$  mapper keeps statistics about
    - (1) the data reaching  $L$
    - (2) the data in left/right subtree under some split  $S$
- **Reducer** aggregates the statistics (1), (2) and determines the best split for each tree node

# PLANET: Components



## ■ Master

- Monitors everything (runs multiple MapReduce jobs)

## ■ Three types of MapReduce jobs:

### ■ (1) MapReduce Initialization (run once first)

- For each attribute identify values to be considered for splits

### ■ (2) MapReduce FindBestSplit (run multiple times)

- MapReduce job to find best split (when there is too much data to fit in memory)

### ■ (3) MapReduce InMemoryBuild (run once last)

- Similar to **BuildSubTree** (but for small data)
- Grows an entire sub-tree once the data fits in memory

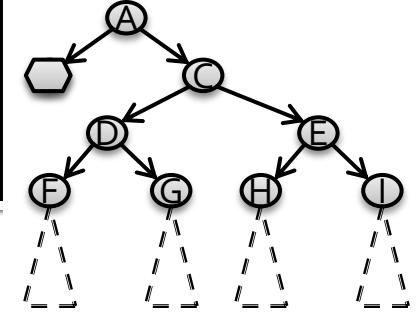
## ■ Model file

- A file describing the state of the model

# PLANET: Components

- (1) Master Node
- (2) MapReduce Initialization (run once first)
- (3) MapReduce FindBestSplit (run multiple times)
- (4) MapReduce InMemoryBuild (run once last)

# PLANET: Master



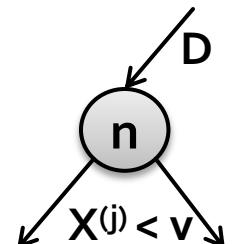
- **Master controls the entire process**
- **Determines the state of the tree and grows it:**
  - (1) Decides if nodes should be split
  - (2) If there is little data entering a tree node, Master runs an InMemoryBuild MapReduce job to grow the entire subtree below that node
  - (3) For larger nodes, Master launches MapReduce FindBestSplit to evaluate candidates for best split
    - Master also collects results from FindBestSplit and chooses the best split for a node
  - (4) Updates the model

# PLANET: Components

- (1) Master Node
- (2) MapReduce Initialization (run once first)
- (3) MapReduce FindBestSplit (run multiple times)
- (4) MapReduce InMemoryBuild (run once last)

# Initialization: Attribute metadata

- **Initialization job:** Identifies all the attribute values which need to be considered for splits
  - Initialization process generates “attribute metadata” to be loaded in memory by other tasks
- **Main question:**  
**Which splits to even consider?**
- **A split is defined by a triple:**  
**(node  $n$ , attribute  $X^{(j)}$ , value  $v$ )**



# Initialization: Attribute metadata

## ■ Which splits to even consider?

- For small data we can sort the values along a particular feature and consider every possible split
- **But data values may not be uniformly populated so many splits may not really make a difference**

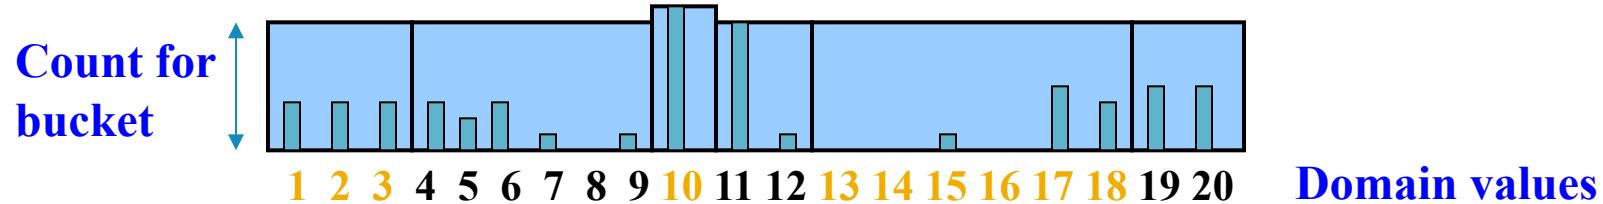
$X^{(i)}$ : 1.2 ↑ 1.3 ↑ 1.4 1.6 2.1 ↑ 7.2 8.1 9.8 10.1 10.2 10.3 10.4 11.5 11.7 12.8 12.9

## ■ Idea: Consider a limited number of splits such that splits “move” about the same amount of data

# Initialization: Attribute metadata

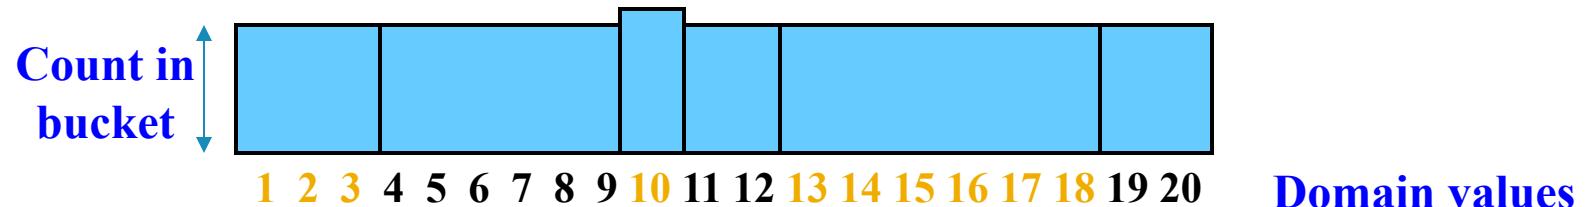
## ■ Splits for numerical attributes:

- For attribute  $X^{(j)}$  we would like to consider every possible value  $v \in O_j$
- Compute an approx. equi-depth histogram on  $D^*$ 
  - Idea: Select buckets such that counts per bucket are equal



- Use boundary points of histogram as splits

# Side note: Computing Equi-Depth



- **Goal:** Equal number of elements per bucket ( $B$  buckets total)
- Construct by first **sorting** and then taking  $B-1$  equally-spaced splits

1 2 2 3 4 7 8 9 10 10 10 10 ↑ 11 11 12 12 ↑ 14 16 16 18 ↑ 19 20 20 20

- **Faster construction:**

Sample & take equally-spaced splits in the sample

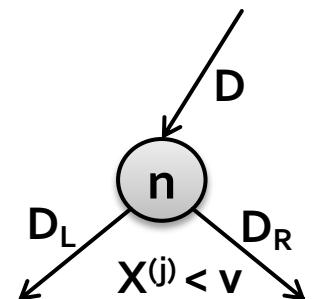
- Nearly equal buckets

# PLANET: Components

- (1) Master Node
- (2) MapReduce Initialization (run once first)
- (3) MapReduce FindBestSplit (run multiple times)
- (4) MapReduce InMemoryBuild (run once last)

# FindBestSplit

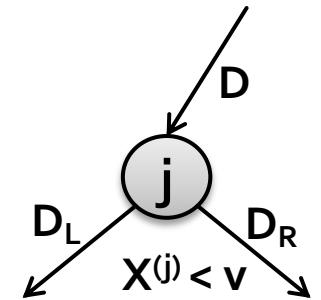
- **Goal:** For a particular split node  $n$  find attribute  $x^{(j)}$  and value  $v$  that **maximize Purity**:
  - $|D| \cdot Var(D) - (|D_L| \cdot Var(D_L) + |D_R| \cdot Var(D_R))$ 
    - $D$  ... training data  $(x_i, y_i)$  reaching the node  $n$
    - $D_L$  ... training data  $x_i$ , where  $x_i^{(j)} < v$
    - $D_R$  ... training data  $x_i$ , where  $x_i^{(j)} \geq v$
  - $Var(D) = \frac{1}{n} \sum_{i \in D} y_i^2 - \left( \frac{1}{n} \sum_{i \in D} y_i \right)^2$



# FindBestSplit

- To compute Purity we need

- $Var(D) = \frac{1}{n} \sum_{i \in D} y_i^2 - \left( \frac{1}{n} \sum_{i \in D} y_i \right)^2$



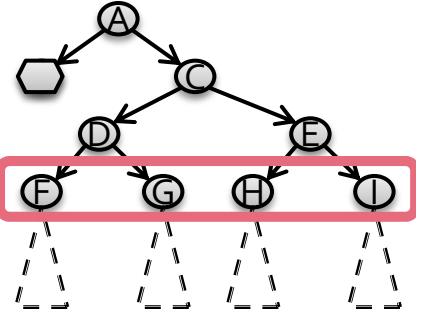
- Important observation: Variance can be computed from sufficient statistics:

$$N, S = \sum y_i, Q = \sum y_i^2$$

- Each Mapper  $m$  processes subset of data  $D_m$ , and computes  $N_m, S_m, Q_m$  for its own  $D_m$
- Reducer combines the statistics and computes global variance and then Purity:

- $Var(D) = \frac{1}{\sum_m N_m} \sum_m Q_m - \left( \frac{1}{\sum_m N_m} \sum_m S_m \right)^2$

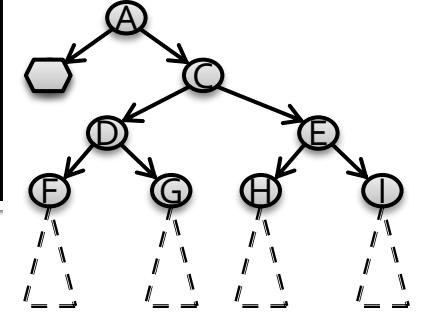
# FindBestSplit: Map



## ■ Mapper:

- Initialized by loading results of **Initialization task**
  - **Current model** (to find which node each datapoint  $x_i$  ends up)
  - **Attribute metadata** (all split points for each attribute)
  - Load the set of **candidate splits**:  $\{(node, attribute, value)\}$
- **For each data record run the Map algorithm:**
  - **For each tree node store statistics of the data entering the node and at the end emit (to all reducers):**
    - $\langle \text{NodeID}, \{ S=\sum y, Q=\sum y^2, N=\sum 1 \} \rangle$
  - **For each split store statistics and at the end emit:**
    - $\langle \text{SplitID}, \{ S, Q, N \} \rangle$ 
      - $\text{SplitID} = (\text{node id}, \text{attribute } X^{(j)}, \text{split value } v)$

# FindBestSplit: Reducer



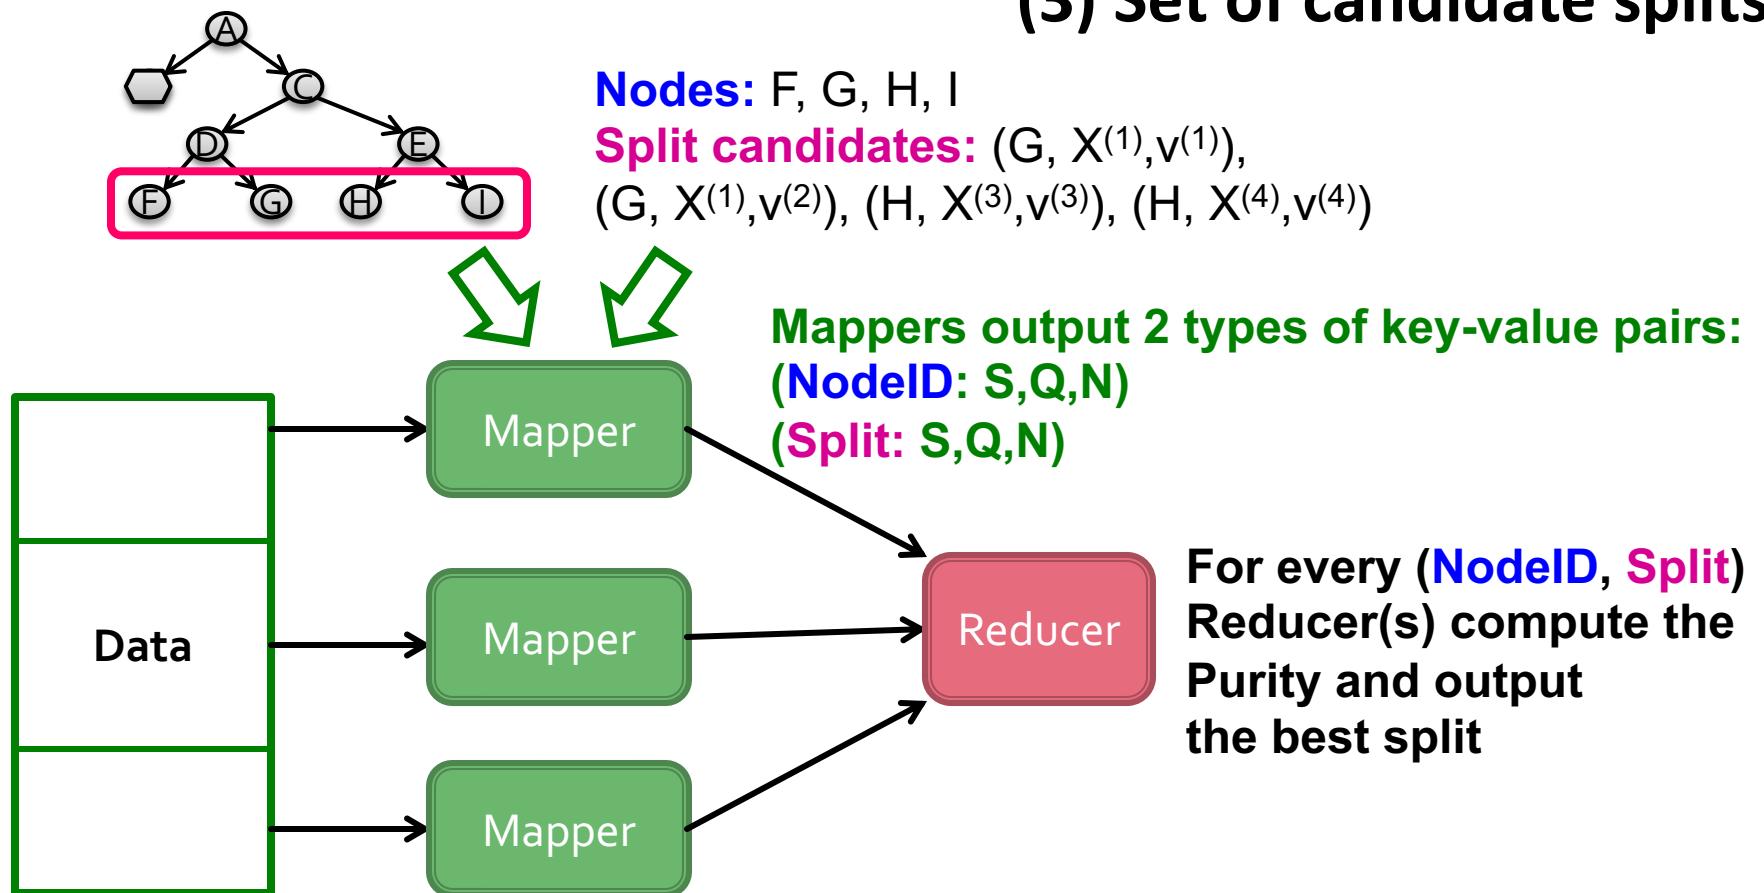
## Reducer:

- (1) Load all the **<NodeID, List {S<sub>m</sub>, Q<sub>m</sub>, N<sub>m</sub>}>** pairs and **aggregate** the per node statistics
- (2) For all the **<SplitID, List {S<sub>m</sub>, Q<sub>m</sub>, N<sub>m</sub>}>** aggregate the statistics

- $Var(D) = \frac{1}{\sum_m N_m} \sum_m Q_m - \left( \frac{1}{\sum_m N_m} \sum_m S_m \right)^2$
- **For each NodeID, output the best split found**

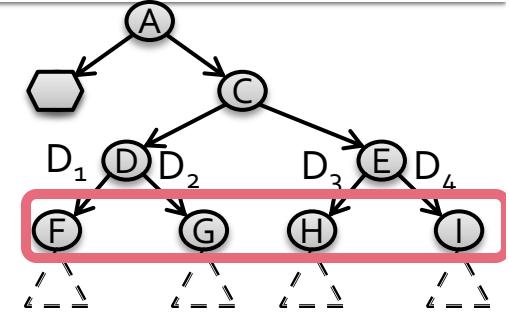
# Overall system architecture

- Master gives the mappers:
  - (1) Tree
  - (2) Set of nodes
  - (3) Set of candidate splits

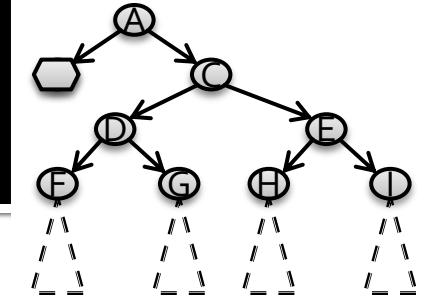


# Overall system architecture

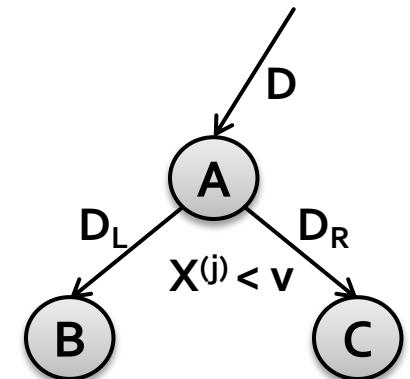
- **Example:** Need to split nodes **F, G, H, I**
- **Map and Reduce:**
  - **FindBestSplit::Map** (each mapper)
    - Load the current model **M**
    - Drop every example  $x_i$  down the tree
    - If it hits **G** or **H**, update in-memory hash tables:
      - For each node:  $T_n: (\text{Node}) \rightarrow \{\mathbf{S}, \mathbf{Q}, \mathbf{N}\}$
      - For each **(Split, Node)**:  $T_{n,j,s}: (\text{Node}, \text{Attribute}, \text{SplitValue}) \rightarrow \{\mathbf{S}, \mathbf{Q}, \mathbf{N}\}$
    - **Map::Finalize**: output the key-value pairs from above hashtables
  - **FindBestSplit::Reduce** (each reducer)
    - Collect:
      - **T1**:  $\langle \text{Node}, \text{List}\{\mathbf{S}, \mathbf{Q}, \mathbf{N}\} \rangle \rightarrow \langle \text{Node}, \{\sum \mathbf{S}, \sum \mathbf{Q}, \sum \mathbf{N}\} \rangle$
      - **T2**:  $\langle (\text{Node}, \text{Attr. Split}), \text{List}\{\mathbf{S}, \mathbf{Q}, \mathbf{N}\} \rangle \rightarrow \langle (\text{Node}, \text{Attr. Split}), \{\sum \mathbf{S}, \sum \mathbf{Q}, \sum \mathbf{N}\} \rangle$
    - Compute Purity for each node using **T1, T2**
    - Return **best split** to Master (which then decides on globally best split)



# Back to the Master



- Collects outputs from **FindBestSplit** reducers  
**<Split.NodeID, Attribute, Value, Purity>**
- For each node decides the best split
  - If data in  $D_L/D_R$  is small enough,  
later run a MapReduce job  
**InMemoryBuild** on the node
  - Else run MapReduce  
**FindBestSplit** job for both  
nodes



# Decision Trees: Conclusion

# Decision Trees

## ■ Characteristics

- Classification & Regression
  - Multiple (~10) classes
- Real valued and categorical features
- Few (hundreds) of features
- Usually dense features
- Complicated decision boundaries
  - Early stopping to avoid overfitting!

## ■ Example applications

- User profile classification
- Landing page bounce prediction

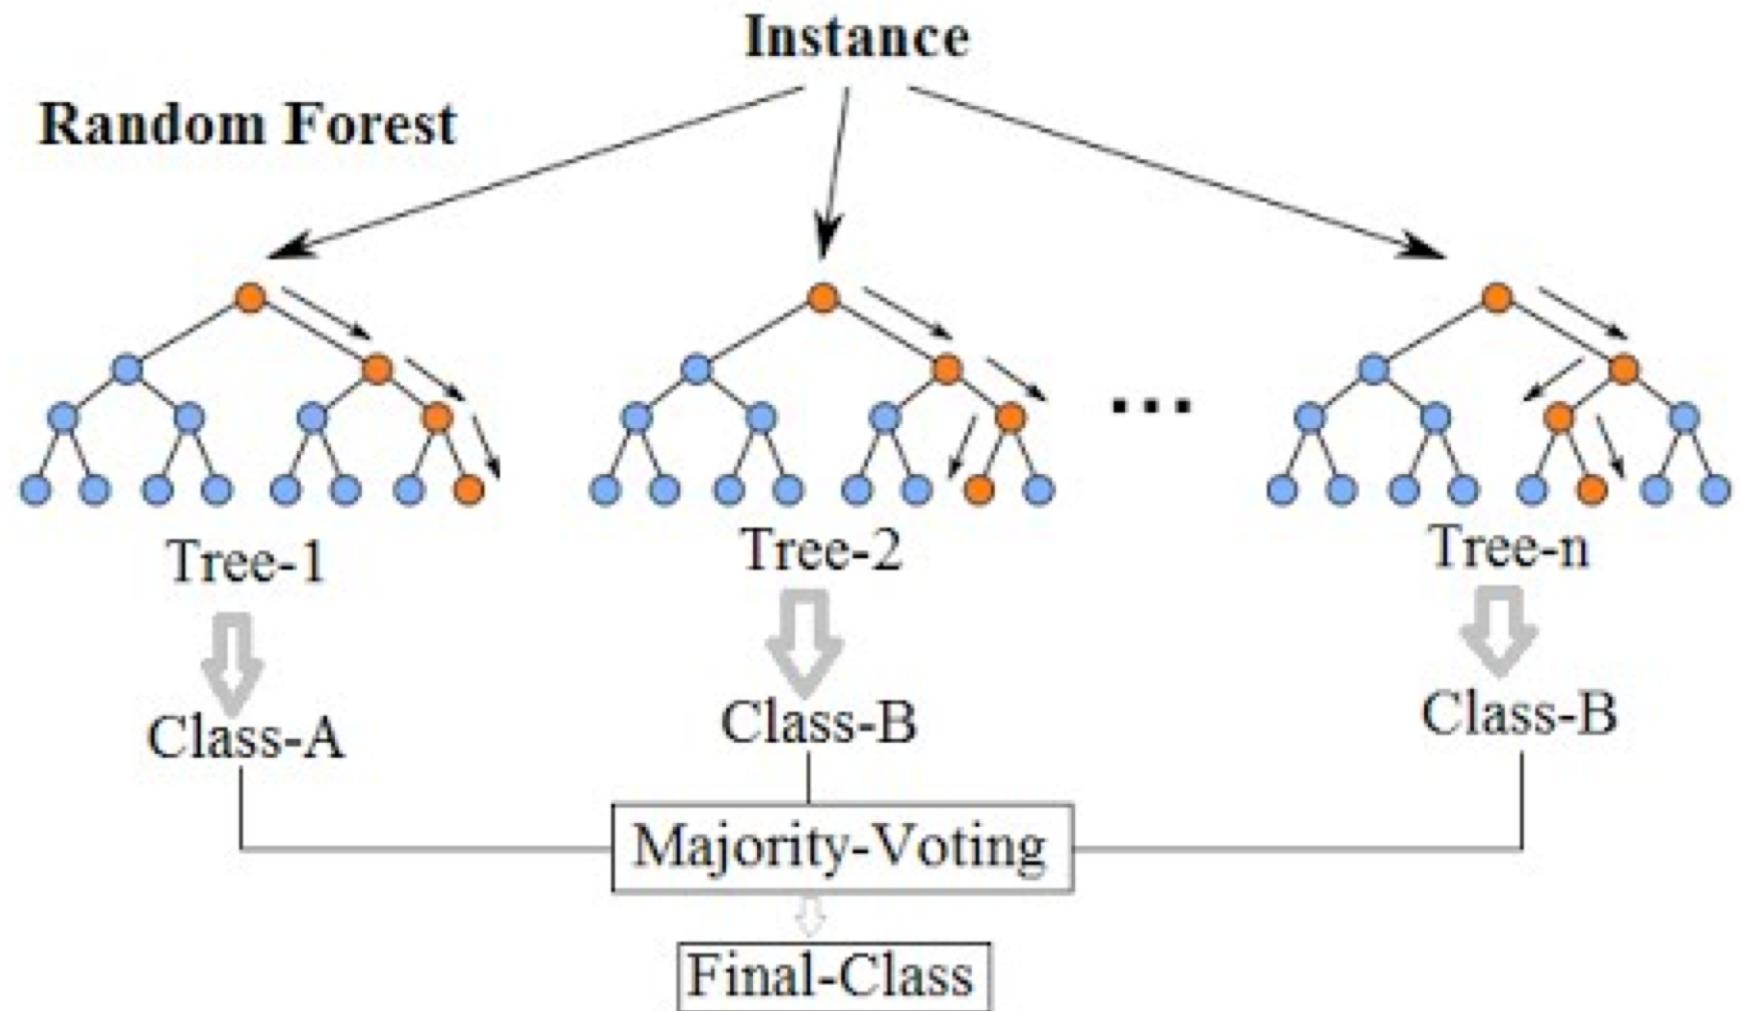
# Decision Trees

- **Decision trees are the single most popular data mining tool:**
  - Easy to understand
  - Easy to implement
  - Easy to use
  - Computationally cheap
  - It's possible to mitigate overfitting (i.e., with ensemble methods)
  - **They do classification as well as regression!**

# Learning Ensembles

- **Learn multiple trees and combine their predictions**
  - Gives better performance in practice
- **Bagging:**
  - Learns multiple trees over independent samples of the training data
    - For a dataset  $D$  on  $n$  data points: Create dataset  $D'$  of  $n$  points but sample from  $D$  with replacement:
      - 33% points in  $D'$  will be duplicates, 66% will be unique
    - Predictions from each tree are averaged to compute the final model prediction

# Bagging Decision Trees



# Bagged Decision Trees

## ■ How to create random samples of $D$ ?

- Compute a hash of a training record's id and tree id
- Use records that hash into a particular range to learn a tree
- This way the same sample is used for all nodes in a tree
- **Note:** This is sampling  $D$  without replacement (but samples of  $D$  should be created with replacement)

# Improvement: Random Forests

- Train a **Bagged Decision Tree**
- But use a modified tree learning algorithm that selects (at each candidate split) **a random subset of the features**
  - If we have  $d$  features, consider  $\sqrt{d}$  random features
- **This is called: Feature bagging**
  - **Benefit:** Breaks correlation between trees
    - If one feature is very strong predictor, then every tree will select it, causing trees to be correlated.
- **Random Forests achieve state-of-the-art results in many classification problems!**