

# Optimizing Submodular Functions

CS246: Mining Massive Datasets  
Jure Leskovec, Stanford University  
<http://cs246.stanford.edu>



# Announcement: Final Exam Logistics

# Final: Logistics

- **Date:**
  - Tuesday, March 19, 3:30-6:30 PM PDT
  - **Location:**
    - if SUNetID[0] in ['a', .. 'l'] then 420-040
    - if SUNetID[0] in ['m', .. 'z'] then Bishop Auditorium
- **Alternate Date:**
  - Monday, March 18, 6:30-9:30 PM PDT
  - **Location:**
    - Gates 104
    - There is still SOME SPACE LEFT!
- **TAs will NOT answer questions during the final**

# Final: SCPD Logistics

You may come to Stanford to take the exam, or...

- **Date:**
  - From Tue, Mar 19, 3:30 PM to Wed, Mar 20, 3:30 PM (all hours in PDT)
  - Agree with your exam monitor on the most convenient 3-hour slot in that window of time
- Exam monitors will receive an email from SCPD with the final exam, which they will in turn forward to you right before the beginning of your 3-hour slot
- Once you completed the exam, make sure to send the file back to your exam monitor (high-quality scanned copy)
- Exam monitors will NOT answer questions during the final

# Final: Instructions

- **Final exam is open book and open notes**
- **A calculator or computer is REQUIRED**
  - You may only use your computer to do arithmetic calculations (i.e., the buttons found on a standard scientific calculator)
  - You may also use your computer to read course notes or the textbook
  - But no Internet/Google/Python access is allowed
- **Practice finals are posted on Piazza!**
- **We recommend bringing a power strip**

# Optimizing Submodular Functions

CS246: Mining Massive Datasets  
Jure Leskovec, Stanford University  
<http://cs246.stanford.edu>



# Recommendations: Diversity

- Redundancy leads to a bad user experience

**Obama Calls for Broad Action on Guns**

**Obama unveils 23 executive actions,  
calls for assault weapons ban**

**Obama seeks assault weapons ban,  
background checks on all gun sales**

- Uncertainty around information need => don't put all eggs in one basket
- How do we optimize for diversity directly?

# Covering the day's news

## France intervenes

# Chuck for Defense

## Argo wins big

## Hagel expects fight

# Monday, January 14, 2013

# Covering the day's news

A word cloud graphic featuring a large oval containing various news headlines and names, such as Hagel, French, Mali, Hollande, Argo, Schumer, and Chuck. The words are colored in shades of red, green, blue, and yellow, and are arranged in a dense, overlapping pattern.

# France intervenes

# Chuck for Defense

# Argo wins big

## New gun proposals

**Monday, January 14, 2013**

# Encode Diversity as Coverage

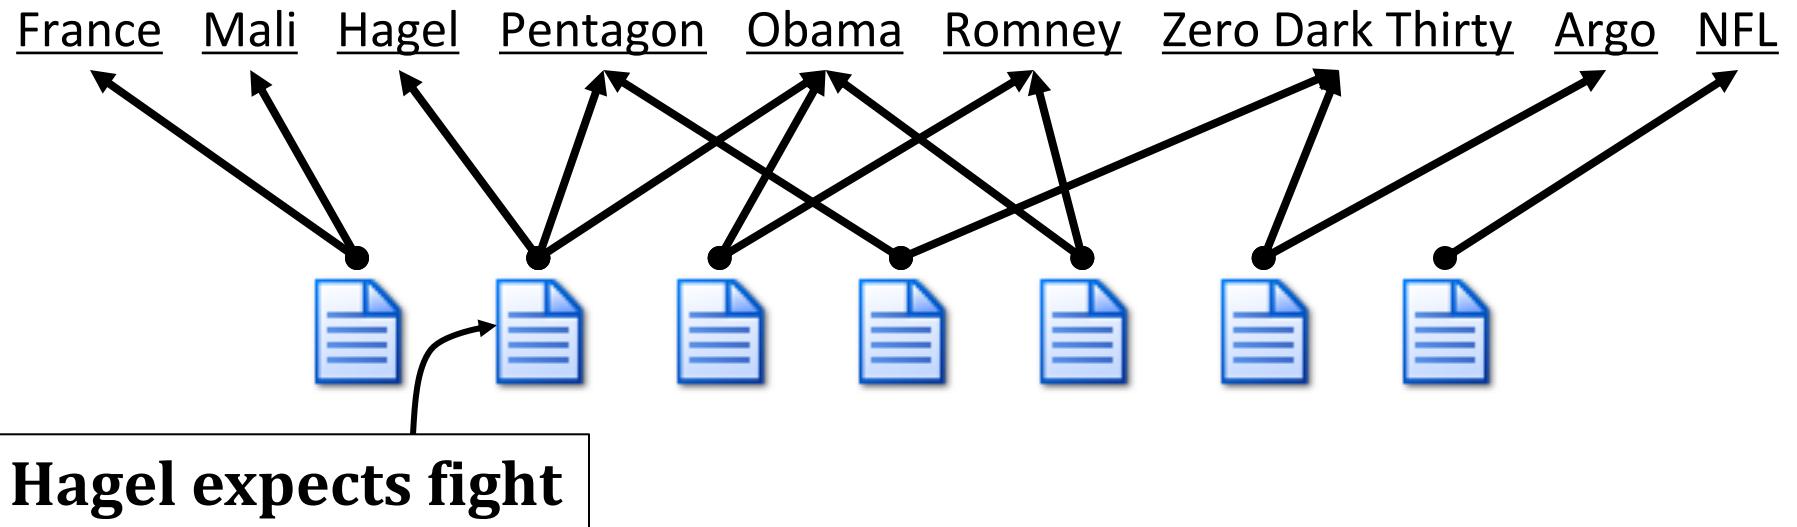
- **Idea:** Encode diversity as coverage problem
  - **Example:** Word cloud of news for a single day
    - Want to select articles so that most words are “covered”



# Diversity as Coverage

# What is being covered?

- Q: What is being covered?
- A: Concepts (In our case: Named entities)



- Q: Who is doing the covering?
- A: Documents

# Simple Abstract Model

- Suppose we are given a set of documents  $D$ 
  - Each document  $d$  covers a set  $X_d$  of words/topics/named entities  $W$
- For a set of documents  $A \subseteq D$  we define

$$F(A) = \left| \bigcup_{i \in A} X_i \right|$$

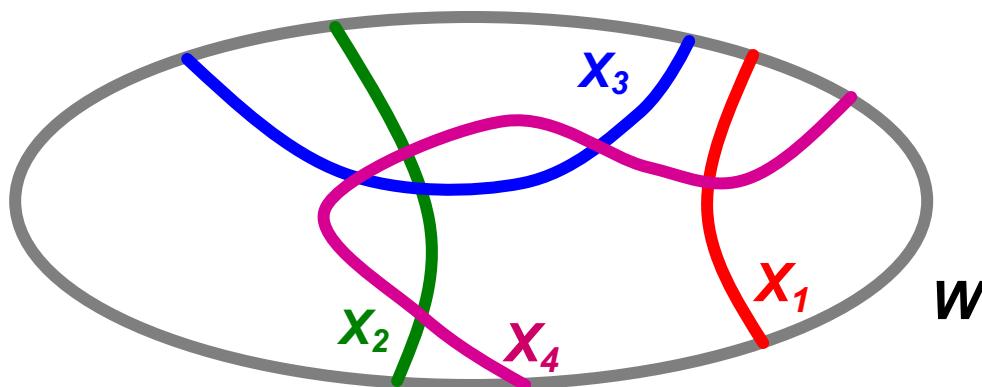
- Goal: We want to

$$\max_{|A| \leq k} F(A)$$

- Note:  $F(A)$  is a set function:  $F(A)$ : Sets  $\rightarrow \mathbb{N}$

# Maximum Coverage Problem

- Given universe of elements  $W = \{w_1, \dots, w_n\}$  and sets  $X_1, \dots, X_m \subseteq W$



- Goal: Find  $k$  sets  $X_i$  that cover the most of  $W$ 
  - More precisely: Find  $k$  sets  $X_i$  whose size of the union is the largest
  - Bad news: A known NP-complete problem

# Simple Greedy Heuristic

## Simple Heuristic: Greedy Algorithm:

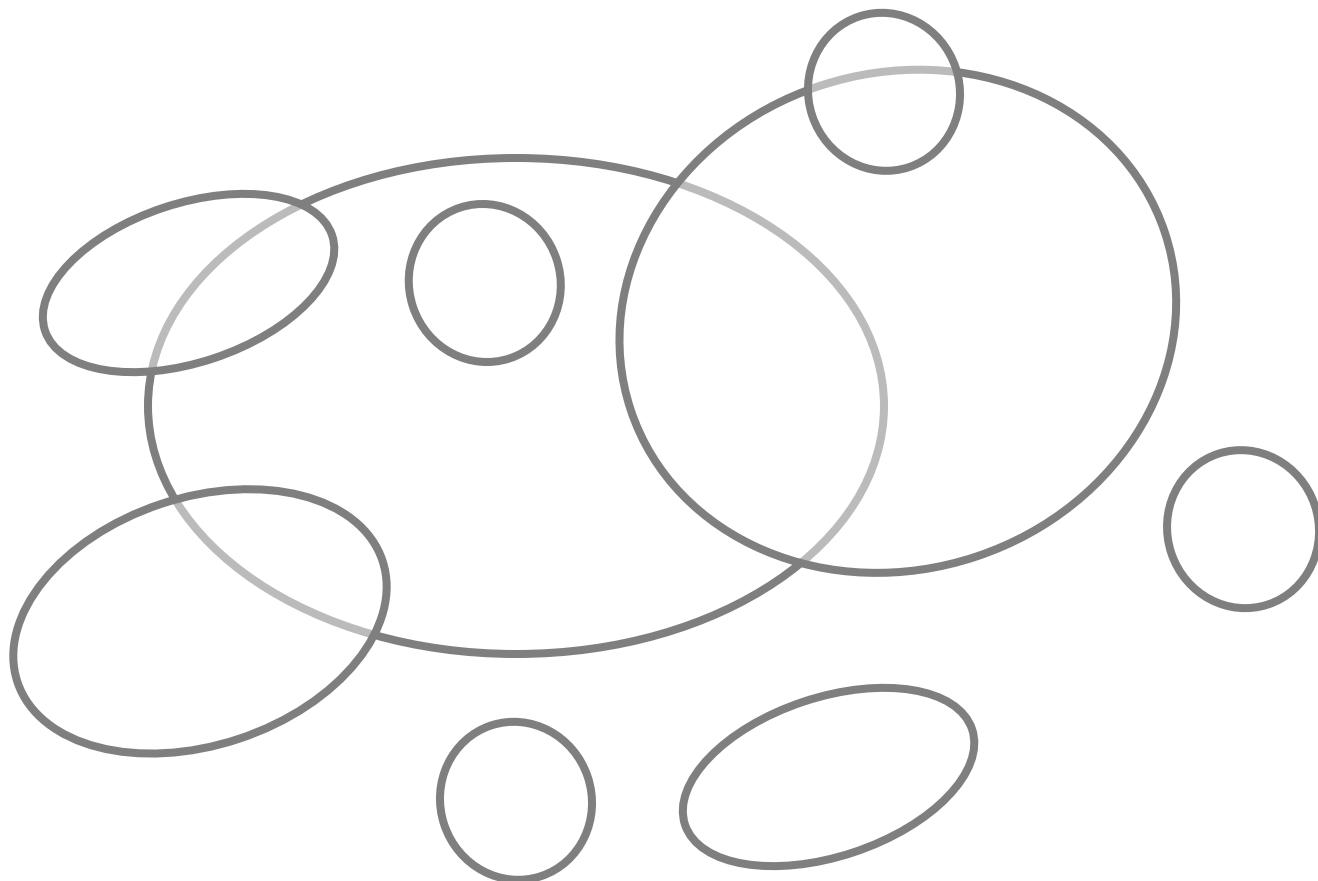
- Start with  $A_0 = \{ \}$
- For  $i = 1 \dots k$ 
  - Find set  $d$  that  $\max F(A_{i-1} \cup \{d\})$
  - Let  $A_i = A_{i-1} \cup \{d\}$

$$F(A) = \left| \bigcup_{d \in A} X_d \right|$$

- Example:
  - Eval.  $F(\{d_1\}), \dots, F(\{d_m\})$ , pick best (say  $d_1$ )
  - Eval.  $F(\{d_1\} \cup \{d_2\}), \dots, F(\{d_1\} \cup \{d_m\})$ , pick best (say  $d_2$ )
  - Eval.  $F(\{d_1, d_2\} \cup \{d_3\}), \dots, F(\{d_1, d_2\} \cup \{d_m\})$ , pick best
  - And so on...

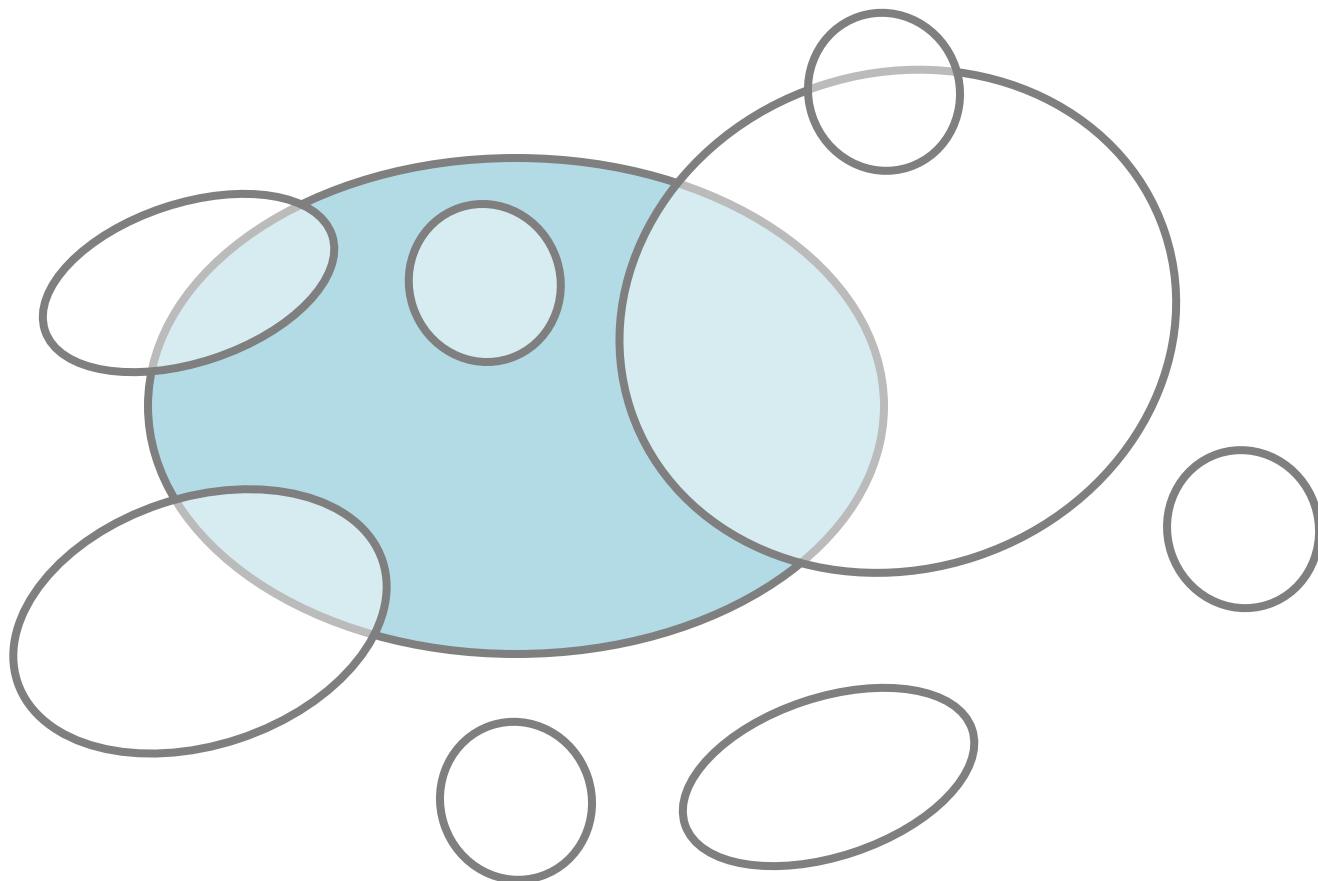
# Simple Greedy Heuristic

- Goal: Maximize the covered area



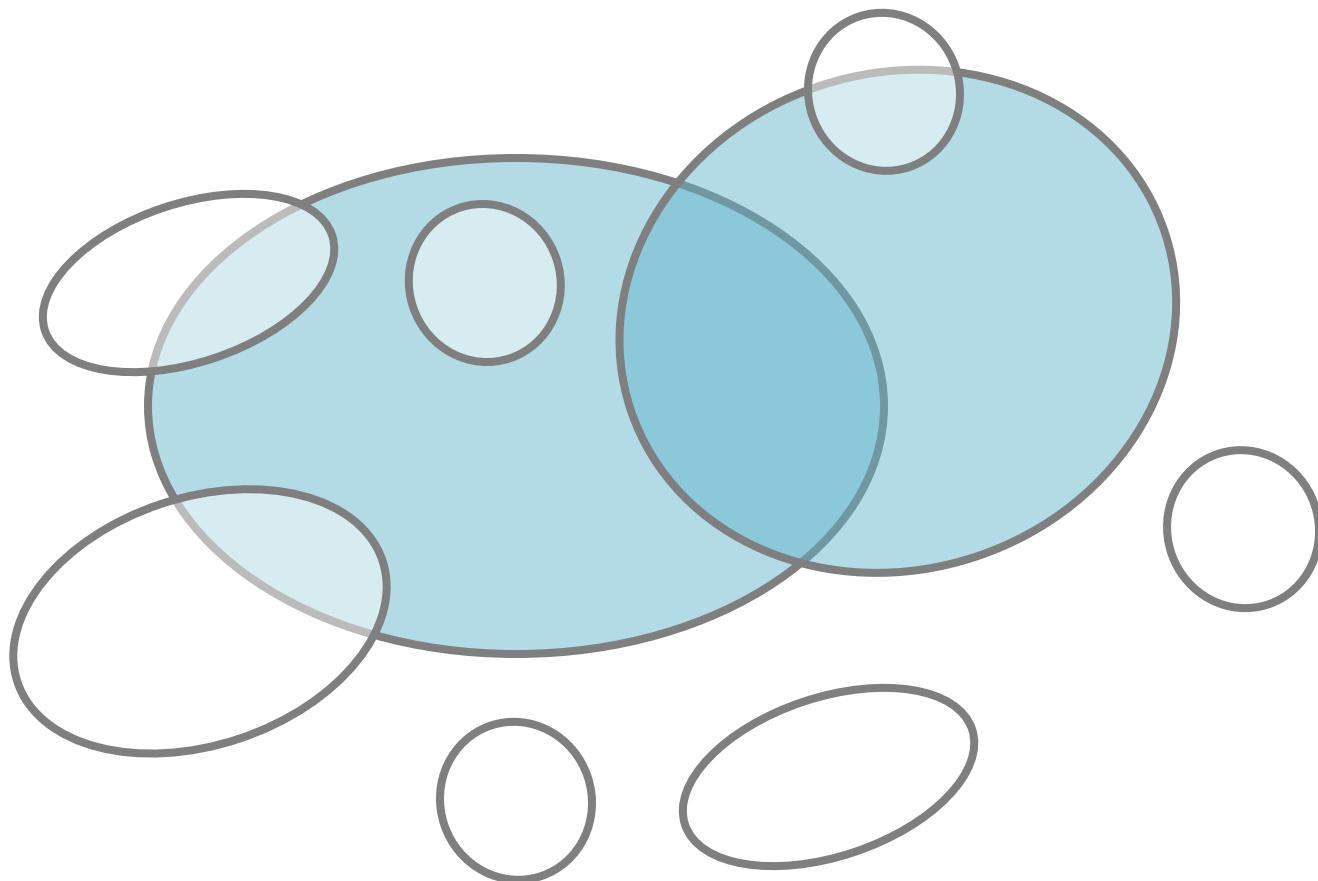
# Simple Greedy Heuristic

- Goal: Maximize the covered area



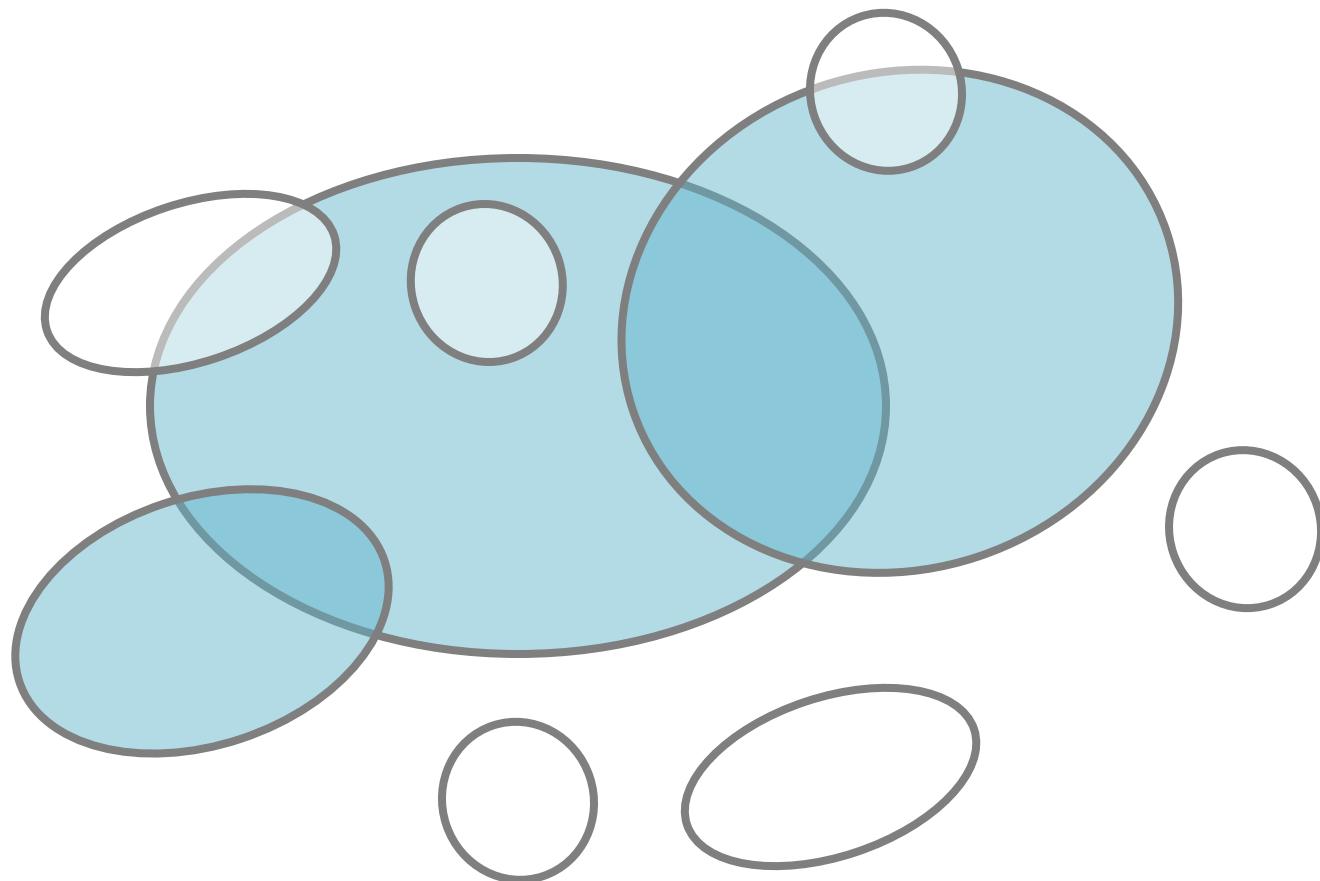
# Simple Greedy Heuristic

- Goal: Maximize the covered area



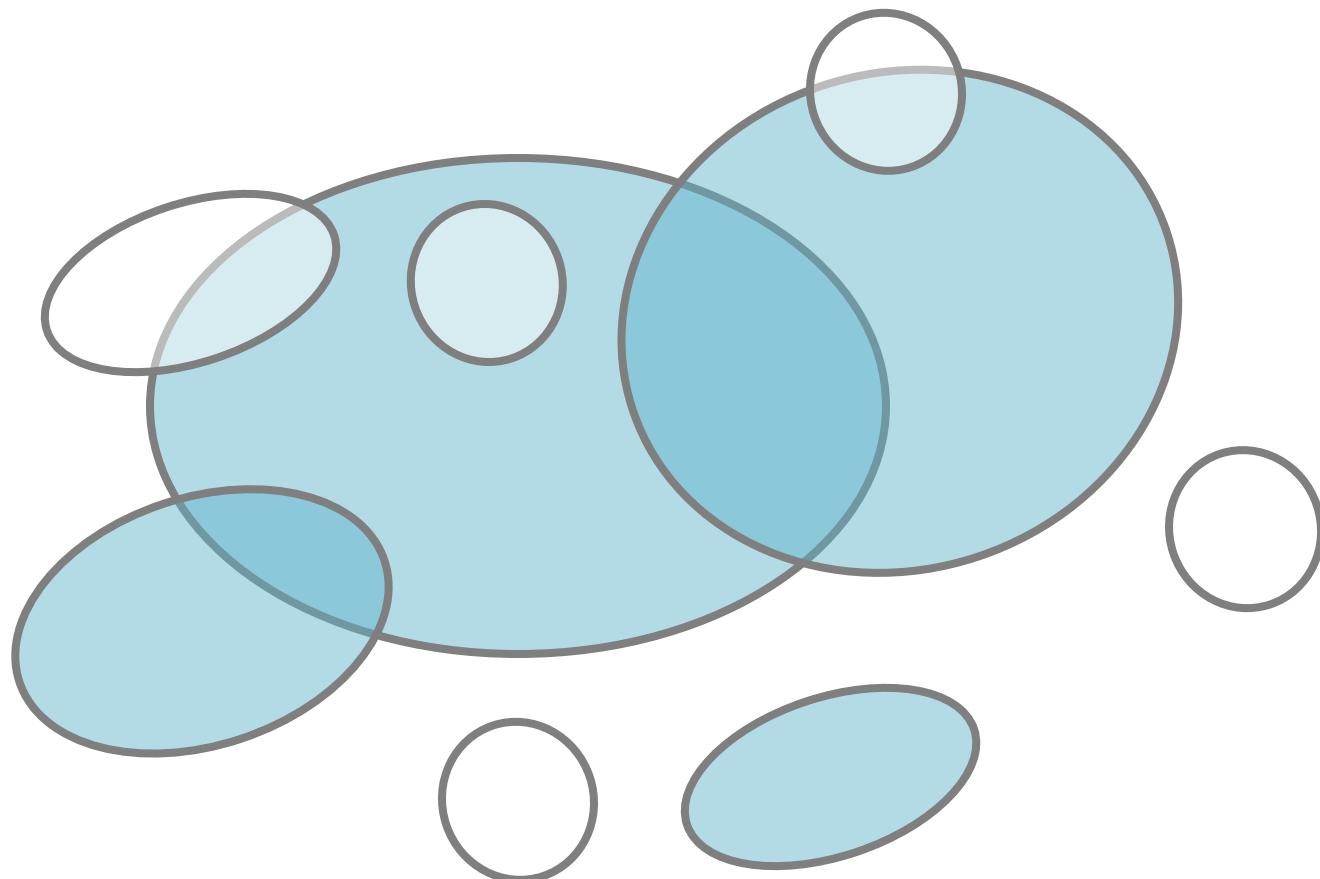
# Simple Greedy Heuristic

- Goal: Maximize the covered area

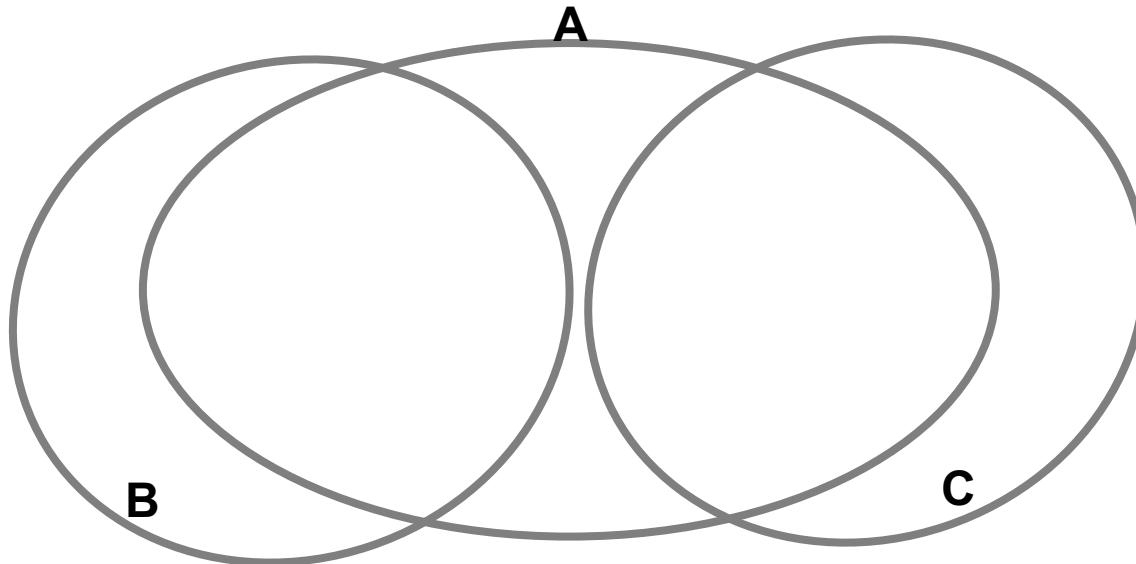


# Simple Greedy Heuristic

- Goal: Maximize the covered area



# When Greedy Heuristic Fails?



- **Goal:** Maximize the size of the covered area
- Greedy first picks A and then C
- But the optimal way would be to pick B and C

# Approximation Guarantee

- **Greedy produces a solution  $A$**   
where:  $F(A) \geq (1-1/e)*OPT$  ( $F(A) \geq 0.63*OPT$ )  
[Nemhauser, Fisher, Wolsey '78]
- **Claim holds for functions  $F(\cdot)$  with 2 properties:**
  - **$F$  is monotone:** (adding more docs doesn't decrease coverage)  
if  $A \subseteq B$  then  $F(A) \leq F(B)$  and  $F(\{\})=0$
  - **$F$  is submodular:**  
adding an element to a set gives less improvement  
than adding it to one of its subsets

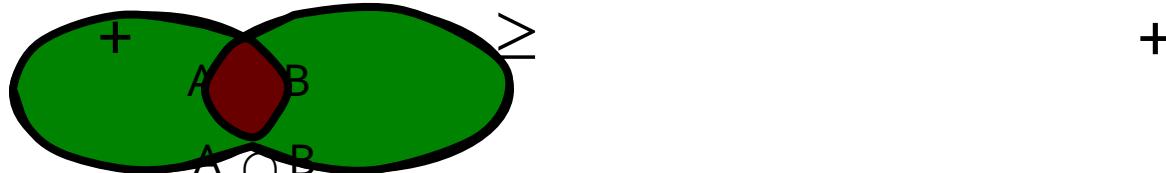
# Submodularity: Definition

## Definition:

- Set function  $F(\cdot)$  is called **submodular** if:

For all  $A, B \subseteq W$ :

$$F(A) + F(B) \geq F(A \cup B) + F(A \cap B)$$



# Submodularity: Or equivalently

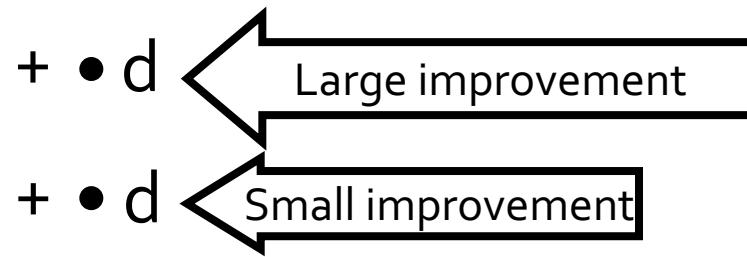
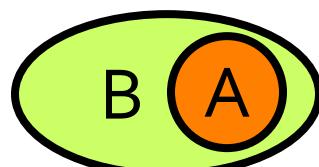
- Diminishing returns characterization

**Equivalent definition:**

- Set function  $F(\cdot)$  is called **submodular** if:

For all  $A \subseteq B$ :

$$\underbrace{F(A \cup \{d\}) - F(A)}_{\text{Gain of adding } d \text{ to a small set}} \geq \underbrace{F(B \cup \{d\}) - F(B)}_{\text{Gain of adding } d \text{ to a large set}}$$



# Example: Set Cover

- $F(\cdot)$  is **submodular**:  $A \subseteq B$

$$F(A \cup \{d\}) - F(A) \geq F(B \cup \{d\}) - F(B)$$

Gain of adding  $d$  to a small set

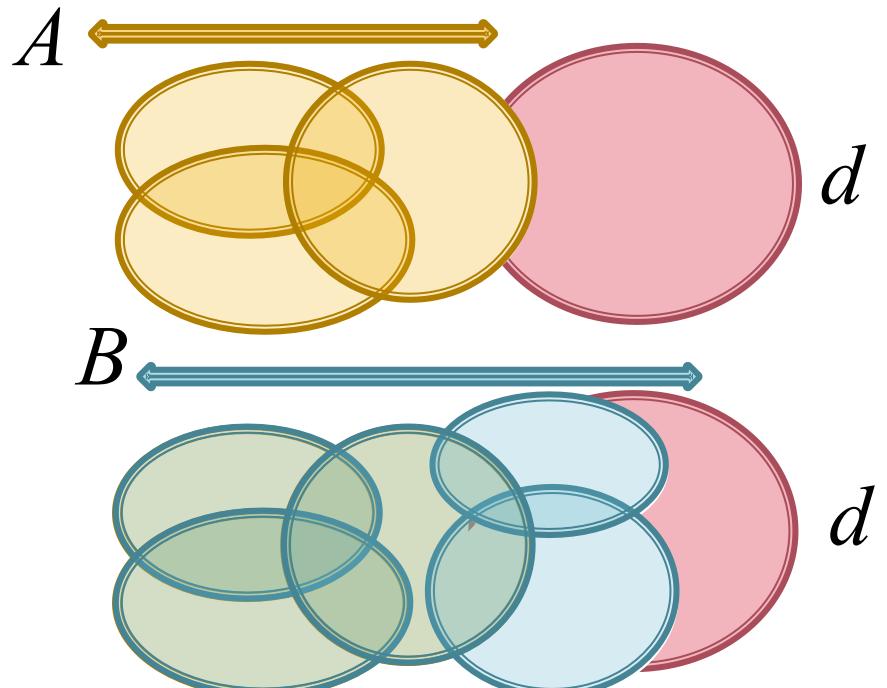
Gain of adding  $d$  to a large set

- **Natural example:**

- Sets  $d_1, \dots, d_m$
- $F(A) = |\bigcup_{i \in A} d_i|$   
(size of the covered area)

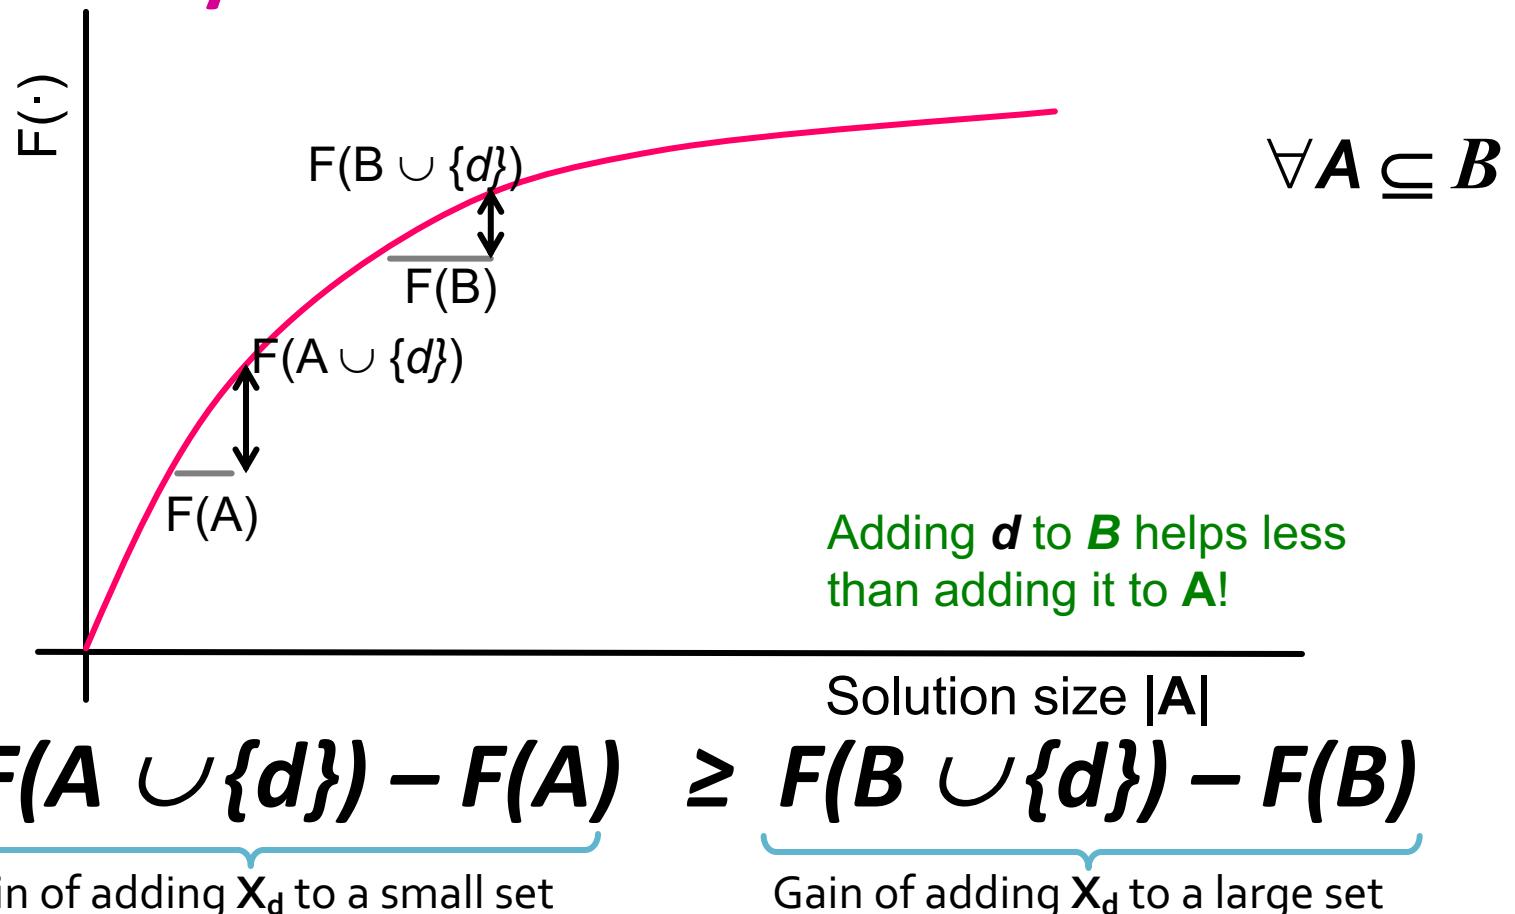
- Claim:

**$F(A)$  is submodular!**



# Submodularity– Diminishing returns

- Submodularity is discrete analogue of concavity



# Submodularity & Concavity

- Marginal gain:

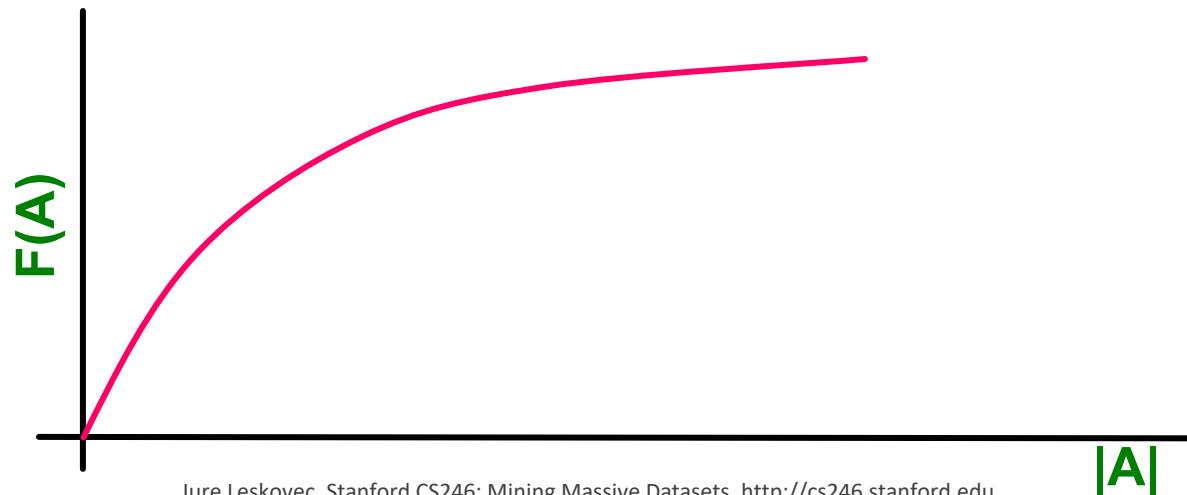
$$\Delta_F(d|A) = F(A \cup \{d\}) - F(A)$$

- Submodular:

$$F(A \cup \{d\}) - F(A) \geq F(B \cup \{d\}) - F(B) \quad A \subseteq B$$

- Concavity:

$$f(a + d) - f(a) \geq f(b + d) - f(b) \quad a \leq b$$

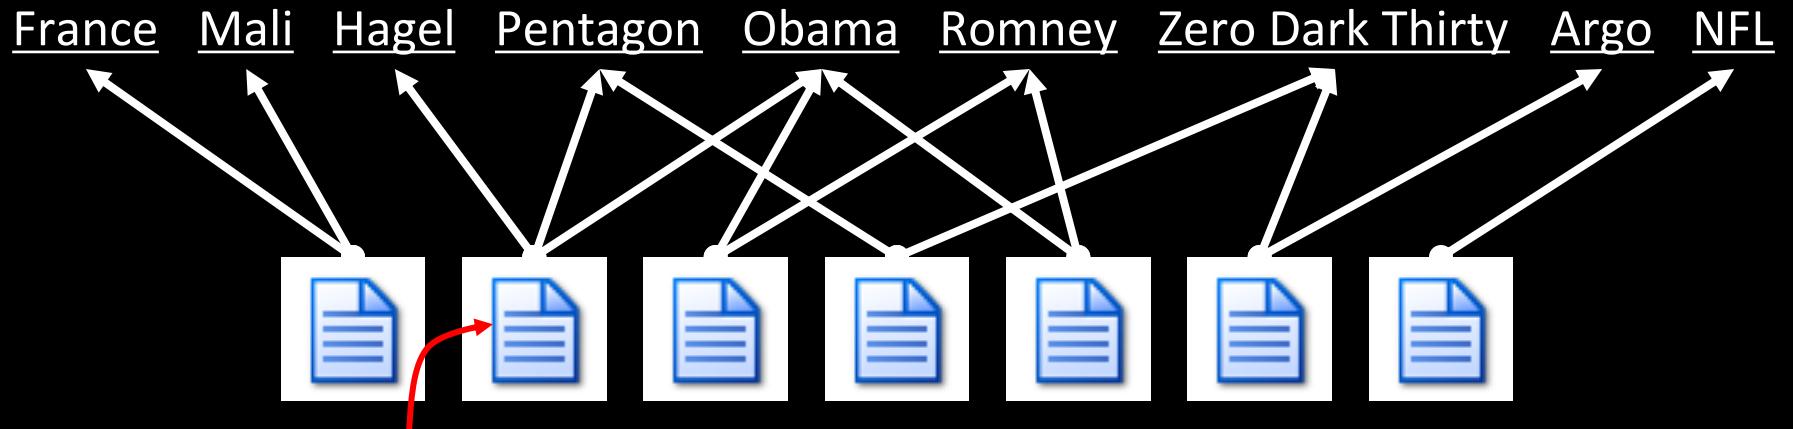


# Submodularity: Useful Fact

- Let  $F_1 \dots F_m$  be **submodular** and  $\lambda_1 \dots \lambda_m > 0$  then  $F(A) = \sum_{i=1}^m \lambda_i F_i(A)$  is **submodular**
  - Submodularity is closed under non-negative linear combinations!
- This is an extremely useful fact:
  - Average of submodular functions is submodular:  
$$F(A) = \sum_i P(i) \cdot F_i(A)$$
  - Multicriterion optimization:  $F(A) = \sum_i \lambda_i F_i(A)$

# Back to our problem

- Q: What is being covered?
- A: Concepts (In our case: Named entities)

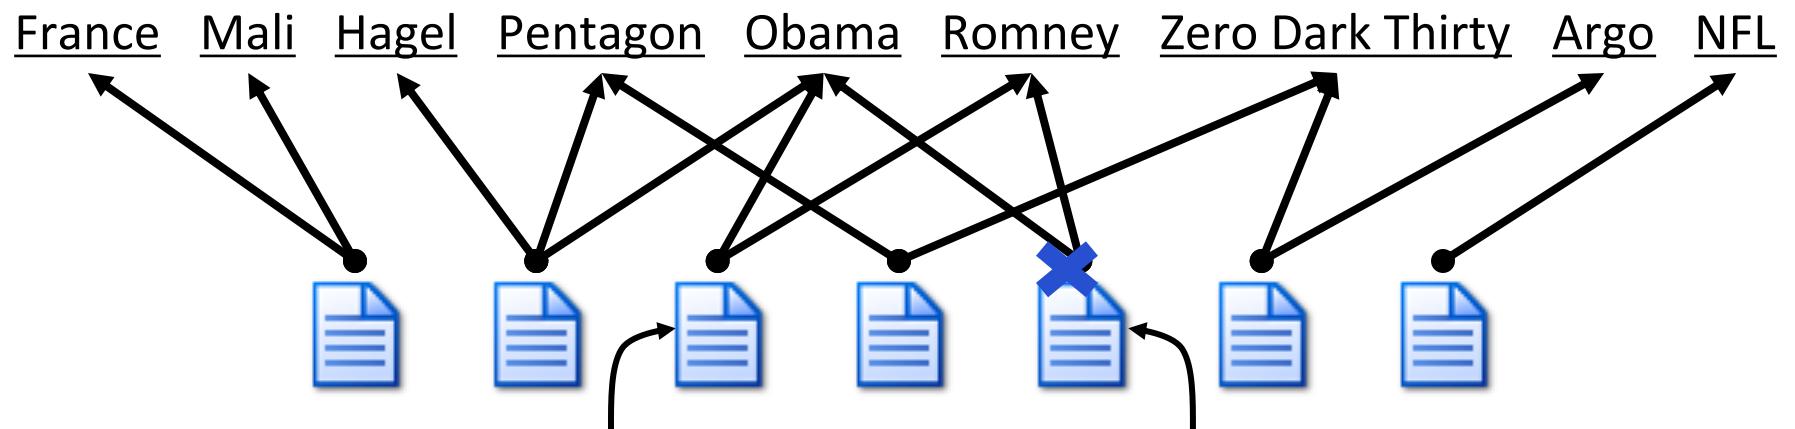


Hagel expects fight

- Q: Who is doing the covering?
- A: Documents

# Back to our Concept Cover Problem

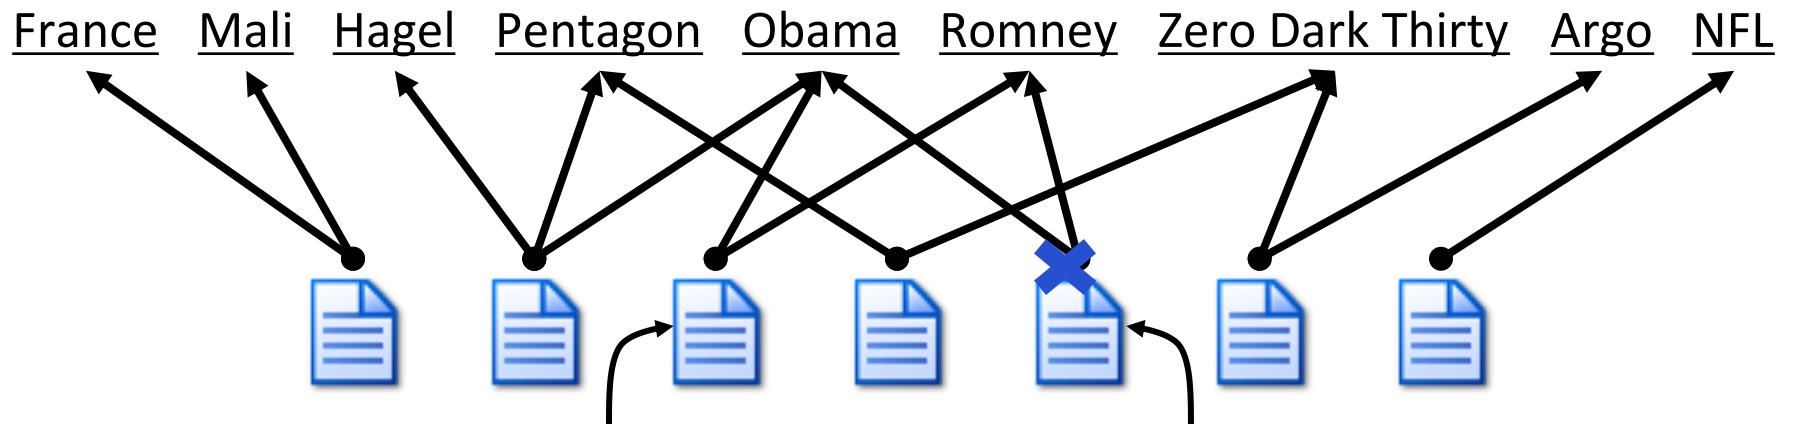
- **Objective:** pick  $k$  docs that cover most concepts



- $F(A)$ : the number of concepts covered by  $A$ 
  - Elements...concepts, Sets ... concepts in docs
  - $F(A)$  is submodular and monotone!
  - We can use **greedy algorithm** to optimize  $F$

# The Set Cover Problem

- **Objective:** pick  $k$  docs that cover most concepts



The good:

Penalizes redundancy

Submodular

The bad:

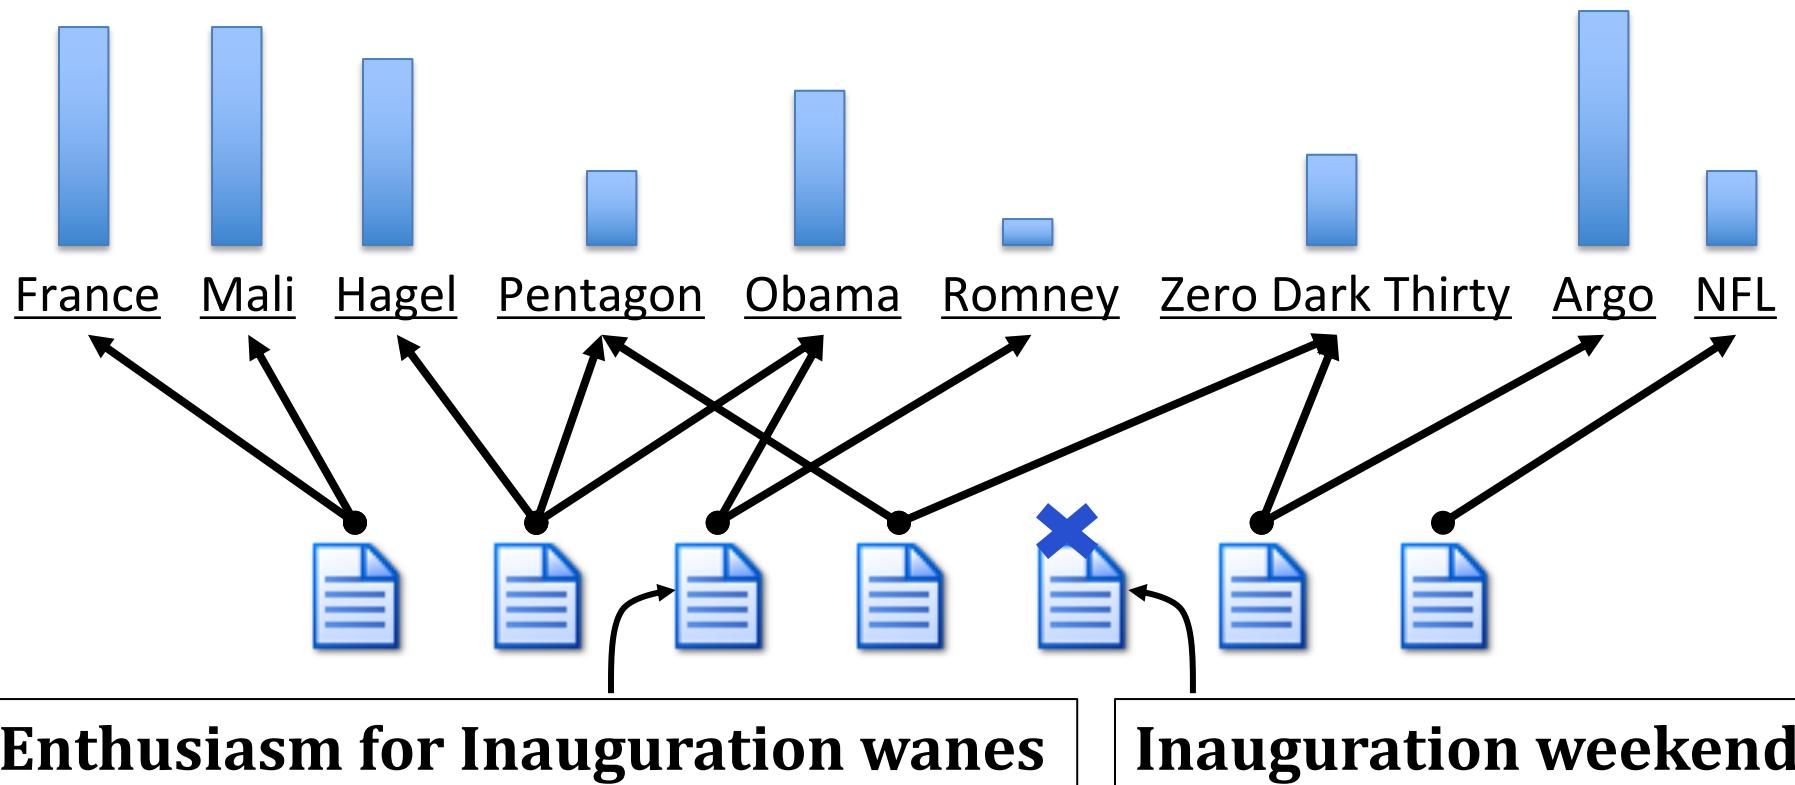
Concept importance?

All-or-nothing too harsh

# Probabilistic Set Cover

# Concept importance?

- **Objective:** pick  $k$  docs that cover most concepts

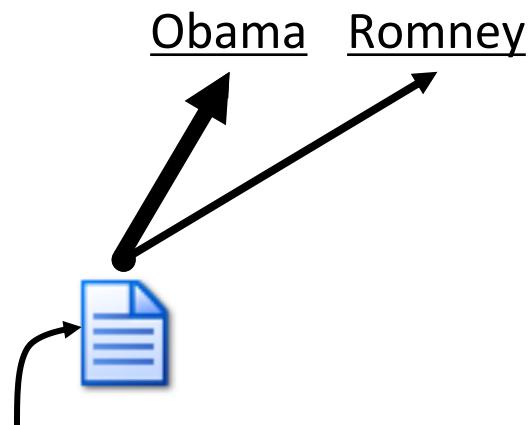


- Each concept  $c$  has importance weight  $w_c$

# All-or-nothing too harsh

- **Document coverage function**

$\text{cover}_d(c) = \text{probability document } d \text{ covers concept } c$   
[e.g., how strongly **d** covers **c**]



Enthusiasm for Inauguration wanes

# Probabilistic Set Cover

- **Document coverage function:**

$\text{cover}_d(c) = \text{probability document } d \text{ covers concept } c$

- $\text{Cover}_d(c)$  can also model how relevant is concept **c** for user **u**

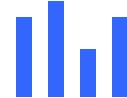
- **Set coverage function:**

$$\text{cover}_{\mathcal{A}}(c) = 1 - \prod_{d \in \mathcal{A}} (1 - \text{cover}_d(c))$$

- Prob. that at least one document in **A** covers **c**

- **Objective:**

$$\max_{\mathcal{A}: |\mathcal{A}| \leq k} F(\mathcal{A}) = \sum_c w_c \text{ cover}_{\mathcal{A}}(c)$$

concept weights 

# Optimizing $F(\mathcal{A})$

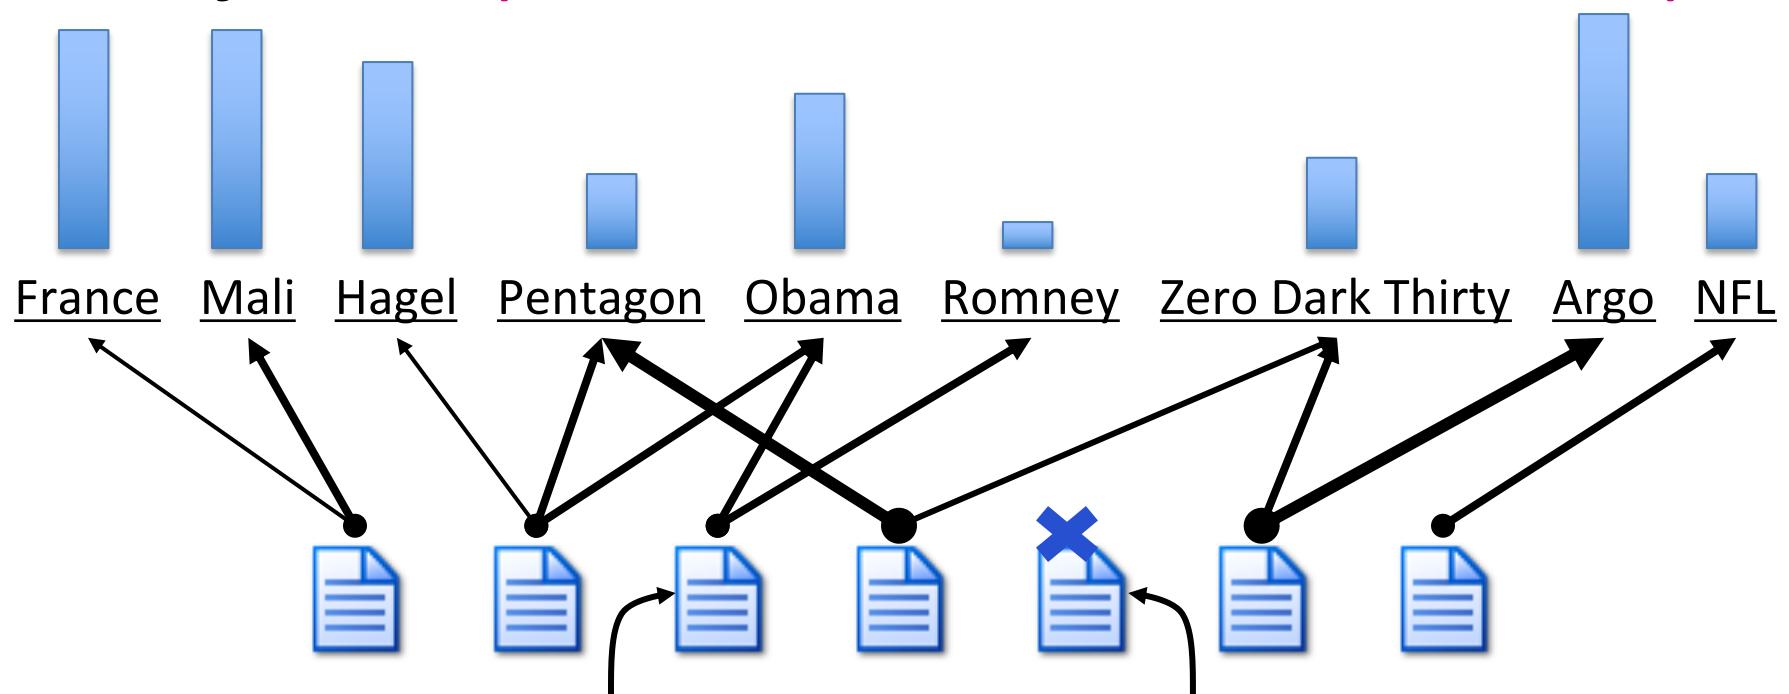
$$\max_{\mathcal{A}: |\mathcal{A}| \leq k} F(\mathcal{A}) = \sum_c w_c \text{ cover}_{\mathcal{A}}(c)$$

- The objective function is also **submodular**
  - Intuitively, it has a **diminishing returns** property
  - Greedy algorithm leads to a  $(1 - 1/e) \sim 63\%$  approximation, i.e., a **near-optimal** solution



# Summary: Probabilistic Set Cover

- **Objective:** pick  $k$  docs that cover most concepts



Enthusiasm for Inauguration wanes

Inauguration weekend

- Each concept  $c$  has importance weight  $w_c$
- Documents partially cover concepts:  $\text{cover}_d(c)$

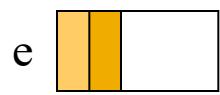
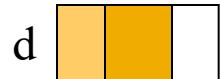
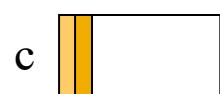
# Lazy Optimization of Submodular Functions

# Submodular Functions

## Greedy

Marginal gain:

$$F(A \cup x) - F(A)$$



Add document with highest marginal gain

### ■ Greedy algorithm is slow!

- At each iteration we need to re-evaluate marginal gains of **all remaining documents**
- Runtime  $O(|D| \cdot K)$  for  **selecting  $K$  documents out of the set of  $D$  of them**

# Speeding up Greedy

- In round  $i$ : So far we have  $A_{i-1} = \{d_1, \dots, d_{i-1}\}$ 
  - Now we pick  $d_i = \arg \max_{d \in V} F(A_{i-1} \cup \{d\}) - F(A_{i-1})$ 
    - Greedy algorithm maximizes the “marginal benefit”  

$$\Delta_i(d) = F(A_{i-1} \cup \{d\}) - F(A_{i-1})$$
- By submodularity property:  

$$F(A_i \cup \{d\}) - F(A_i) \geq F(A_j \cup \{d\}) - F(A_j)$$
 for  $i < j$
- Observation: By submodularity:  
For every  $d \in D$   
 $\Delta_i(d) \geq \Delta_j(d)$  for  $i < j$  since  $A_i \subseteq A_j$   

$$\Delta_i(d) \geq \Delta_j(d)$$
- Marginal benefits  $\Delta_i(d)$  only shrink!  
(as  $i$  grows)

Selecting document  $d$  in step  $i$  covers more words than selecting  $d$  at step  $j$  ( $j > i$ )



# Lazy Greedy

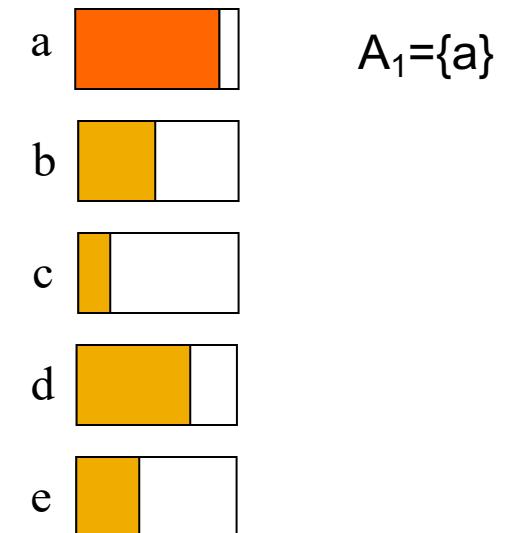
- **Idea:**

- Use  $\Delta_i$  as upper-bound on  $\Delta_j$  ( $j > i$ )

- **Lazy Greedy:**

- Keep an ordered list of marginal benefits  $\Delta_i$  from previous iteration
- Re-evaluate  $\Delta_i$  **only** for top element
- Re-sort and prune

(Upper bound on)  
Marginal gain  $\Delta_1$



$$F(A \cup \{d\}) - F(A) \geq F(B \cup \{d\}) - F(B) \quad A \subseteq B$$

# Lazy Greedy

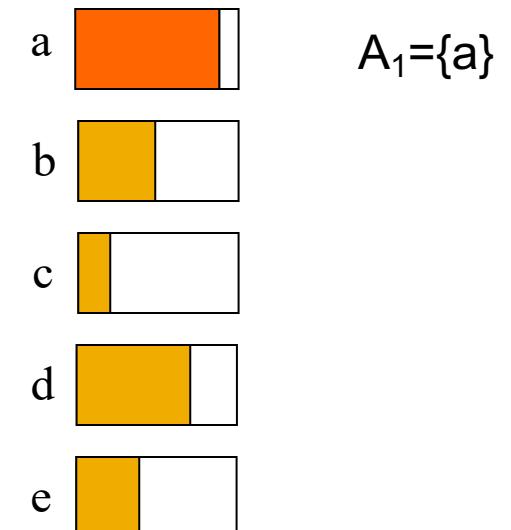
- **Idea:**

- Use  $\Delta_i$  as upper-bound on  $\Delta_j$  ( $j > i$ )

- **Lazy Greedy:**

- Keep an ordered list of marginal benefits  $\Delta_i$  from previous iteration
- Re-evaluate  $\Delta_i$  **only** for top element
- Re-sort and prune

Upper bound on Marginal gain  $\Delta_2$



$$F(A \cup \{d\}) - F(A) \geq F(B \cup \{d\}) - F(B) \quad A \subseteq B$$

# Lazy Greedy

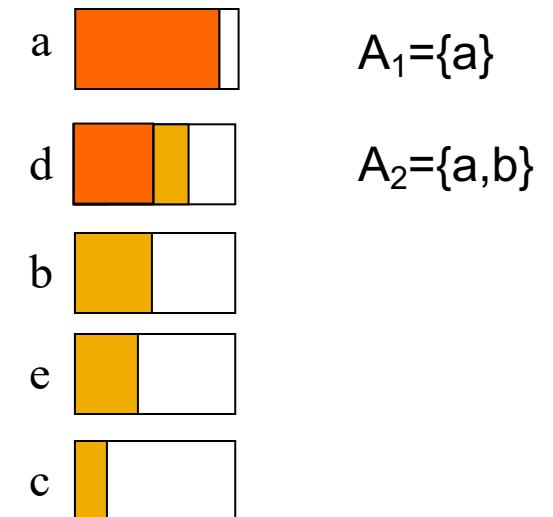
- **Idea:**

- Use  $\Delta_i$  as upper-bound on  $\Delta_j$  ( $j > i$ )

- **Lazy Greedy:**

- Keep an ordered list of marginal benefits  $\Delta_i$  from previous iteration
- Re-evaluate  $\Delta_i$  **only** for top element
- Re-sort and prune

Upper bound on  
Marginal gain  $\Delta_2$

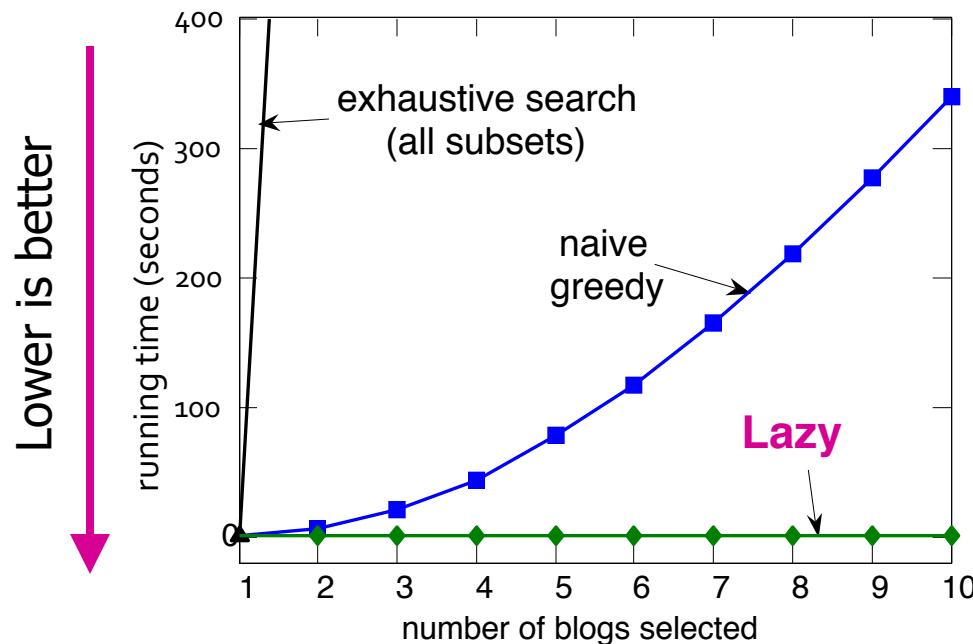


$$F(A \cup \{d\}) - F(A) \geq F(B \cup \{d\}) - F(B) \quad A \subseteq B$$

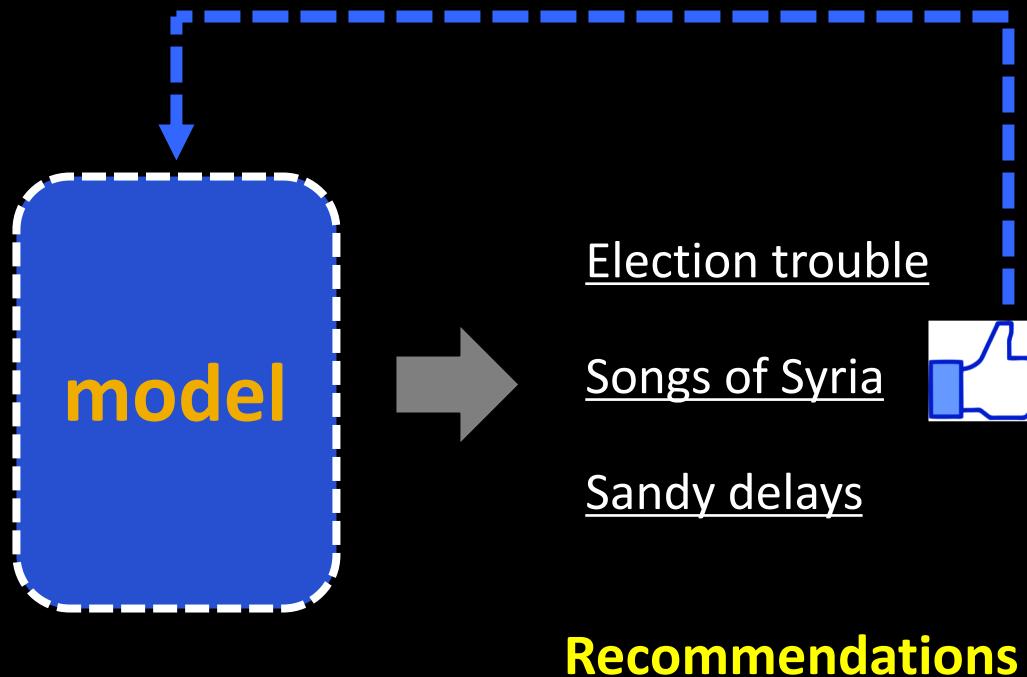
# Summary so far

## ■ Summary so far:

- Diversity can be formulated as a set cover
- Set cover is submodular optimization problem
- Can be (approximately) solved using greedy algorithm
- Lazy-greedy gives significant speedup

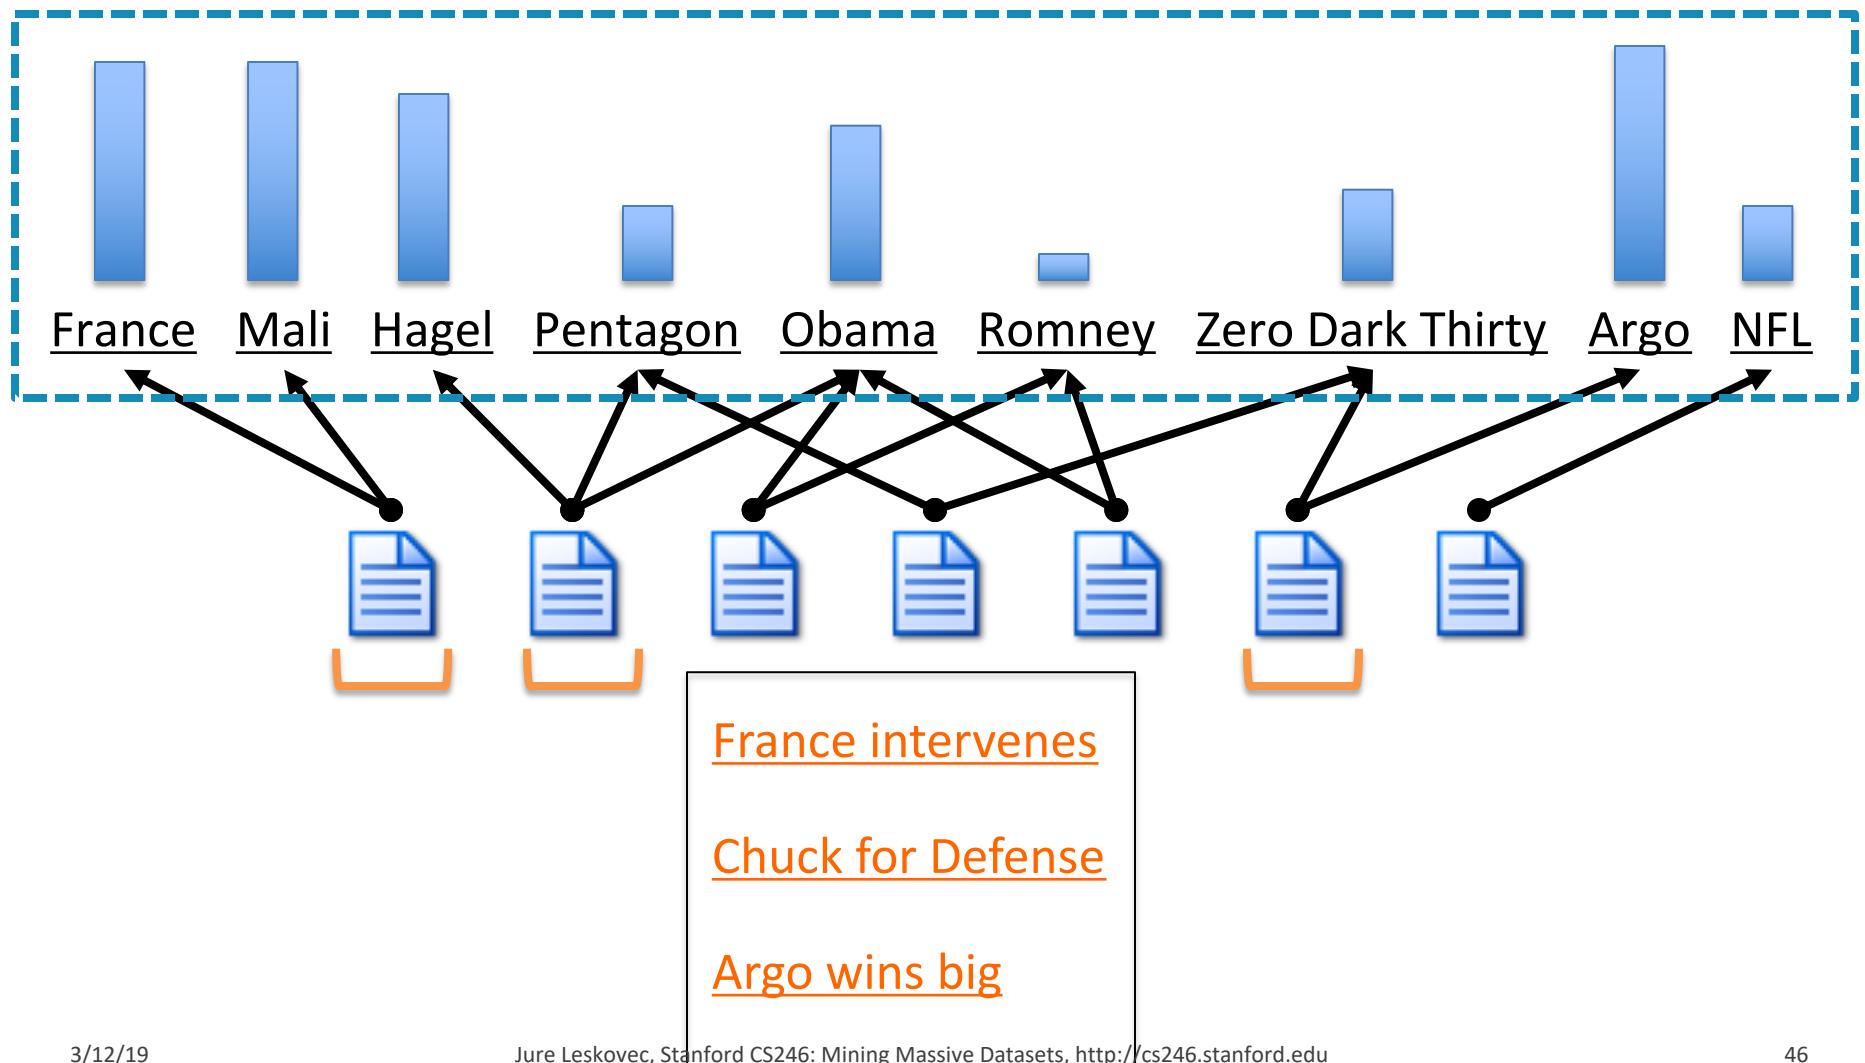


# But what about personalization?



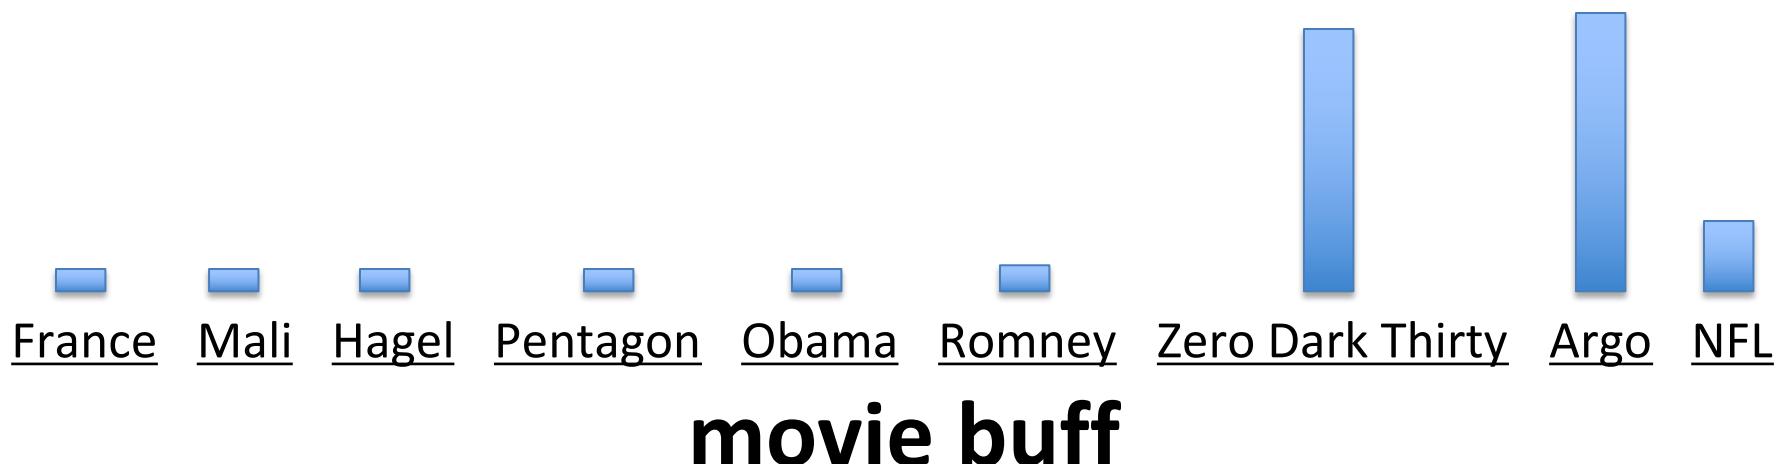
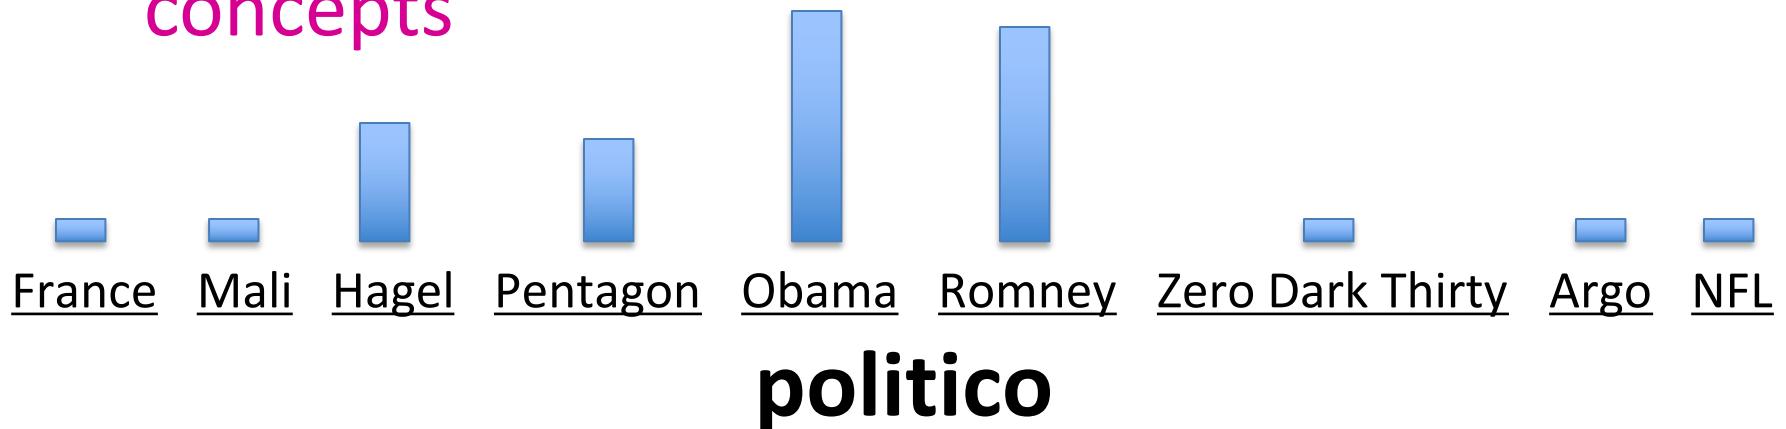
# Concept Coverage

We assumed same concept weighting for all users



# Personal Concept Weights

- Each user has **different** preferences over concepts



# Personal concept weights

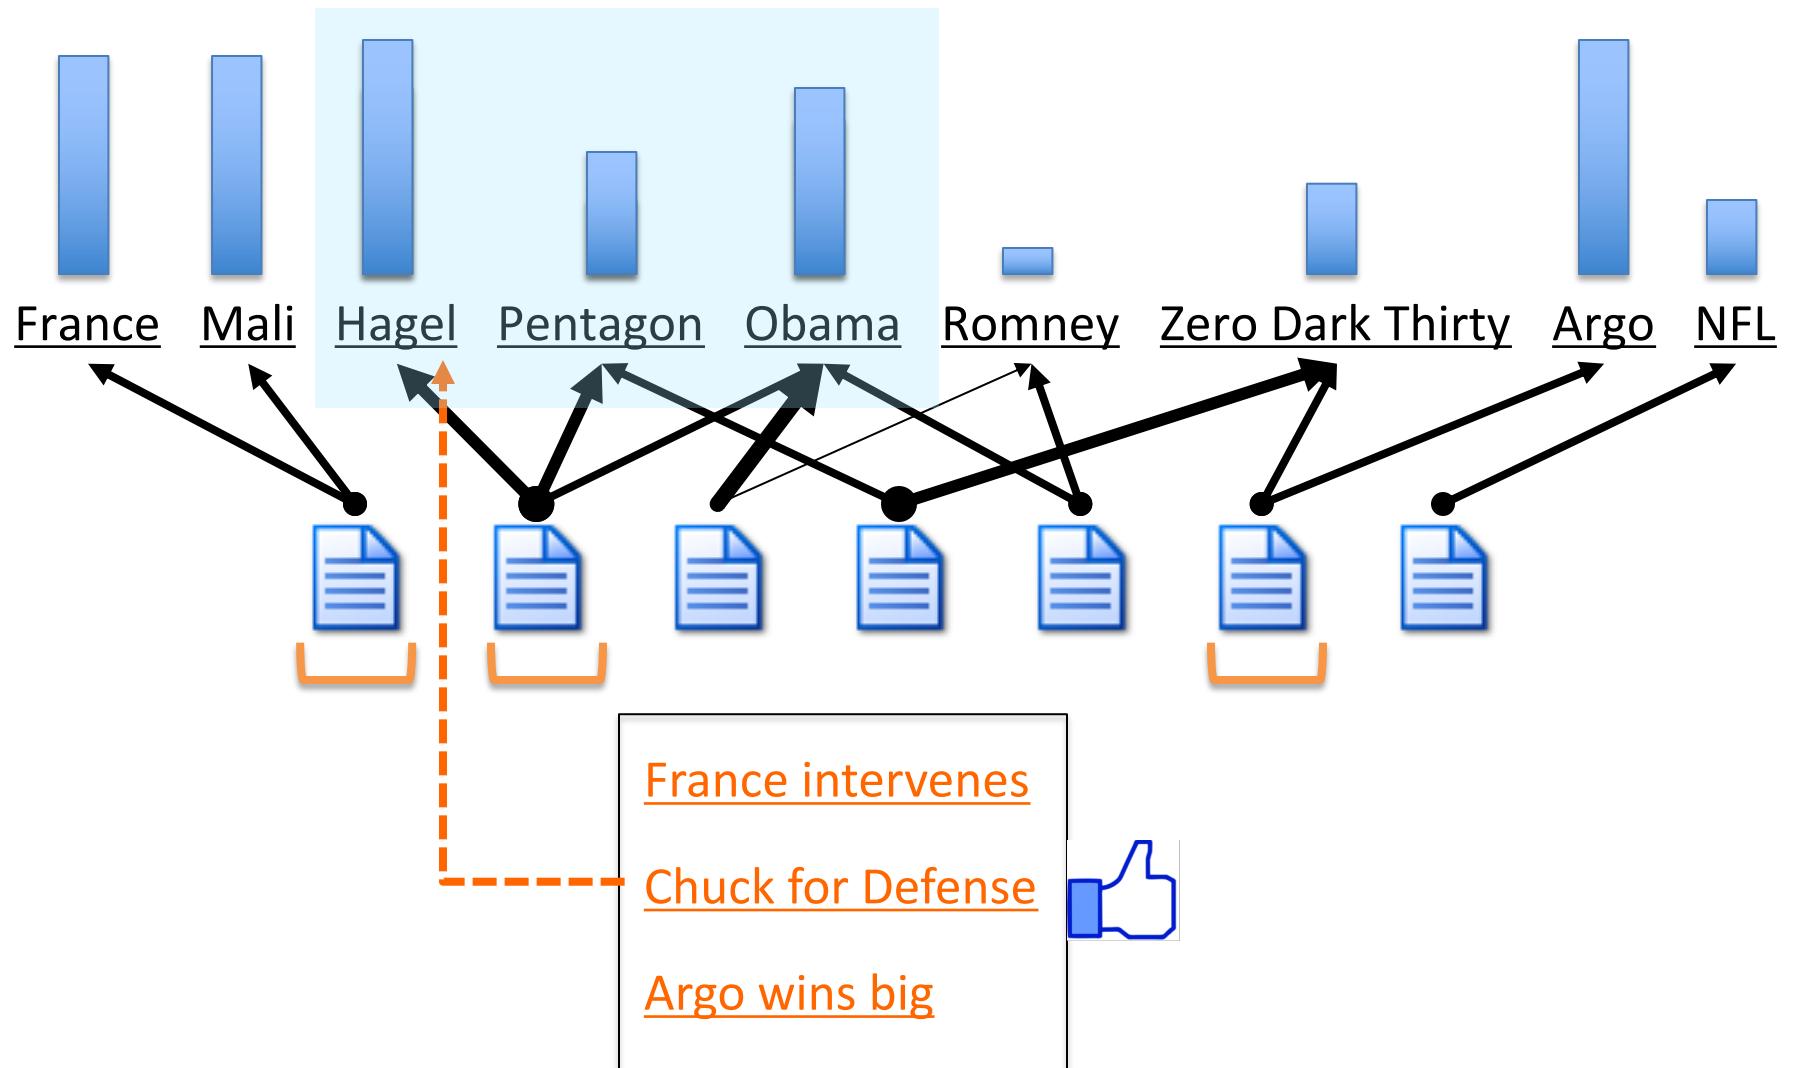
- Assume each user  $u$  has **different** preference vector  $\mathbf{w}_c^{(u)}$  over concepts  $c$

$$\max_{\mathcal{A}:|\mathcal{A}| \leq k} F(\mathcal{A}) = \sum_c w_c \text{ cover}_{\mathcal{A}}(c)$$


$$\max_{\mathcal{A}:|\mathcal{A}| \leq k} F(\mathcal{A}) = \sum_c w_c^{(u)} \text{ cover}_{\mathcal{A}}(c)$$

- Goal:** Learn personal concept weights from user feedback

# Interactive Concept Coverage



# Multiplicative Weights (MW)

## ■ Multiplicative Weights algorithm

- Assume each concept  $c$  has weight  $w_c$
- We recommend document  $d$  and receive feedback, say  $r = +1$  or  $-1$
- **Update the weights:**
  - For each  $c \in X_d$  set  $w_c = \beta^r w_c$ 
    - If concept  $c$  appears in doc  $d$  and we received positive feedback  $r=+1$  then we increase the weight  $w_c$  by multiplying it by  $\beta$  ( $\beta > 1$ ) otherwise we decrease the weight (divide by  $\beta$ )
  - Normalize weights so that  $\sum_c w_c = 1$

# Summary of the Algorithm

## ■ Steps of the algorithm:

1. Identify **items** to recommend from
2. Identify **concepts** [what makes items redundant?]
3. **Weigh** concepts by general importance
4. Define **item-concept coverage function**
5. **Select** items using probabilistic set cover
6. Obtain **feedback, update weights**